

# CLOUD NATIVE CERTIFIED

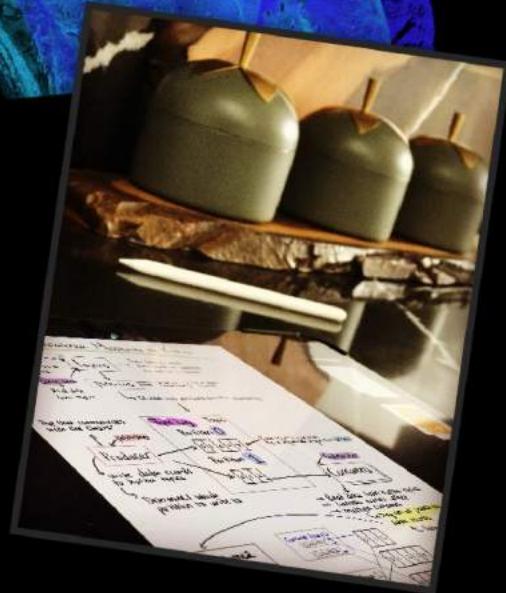


# KUBERNETES ADMINISTRATOR

ADNAN RASHID

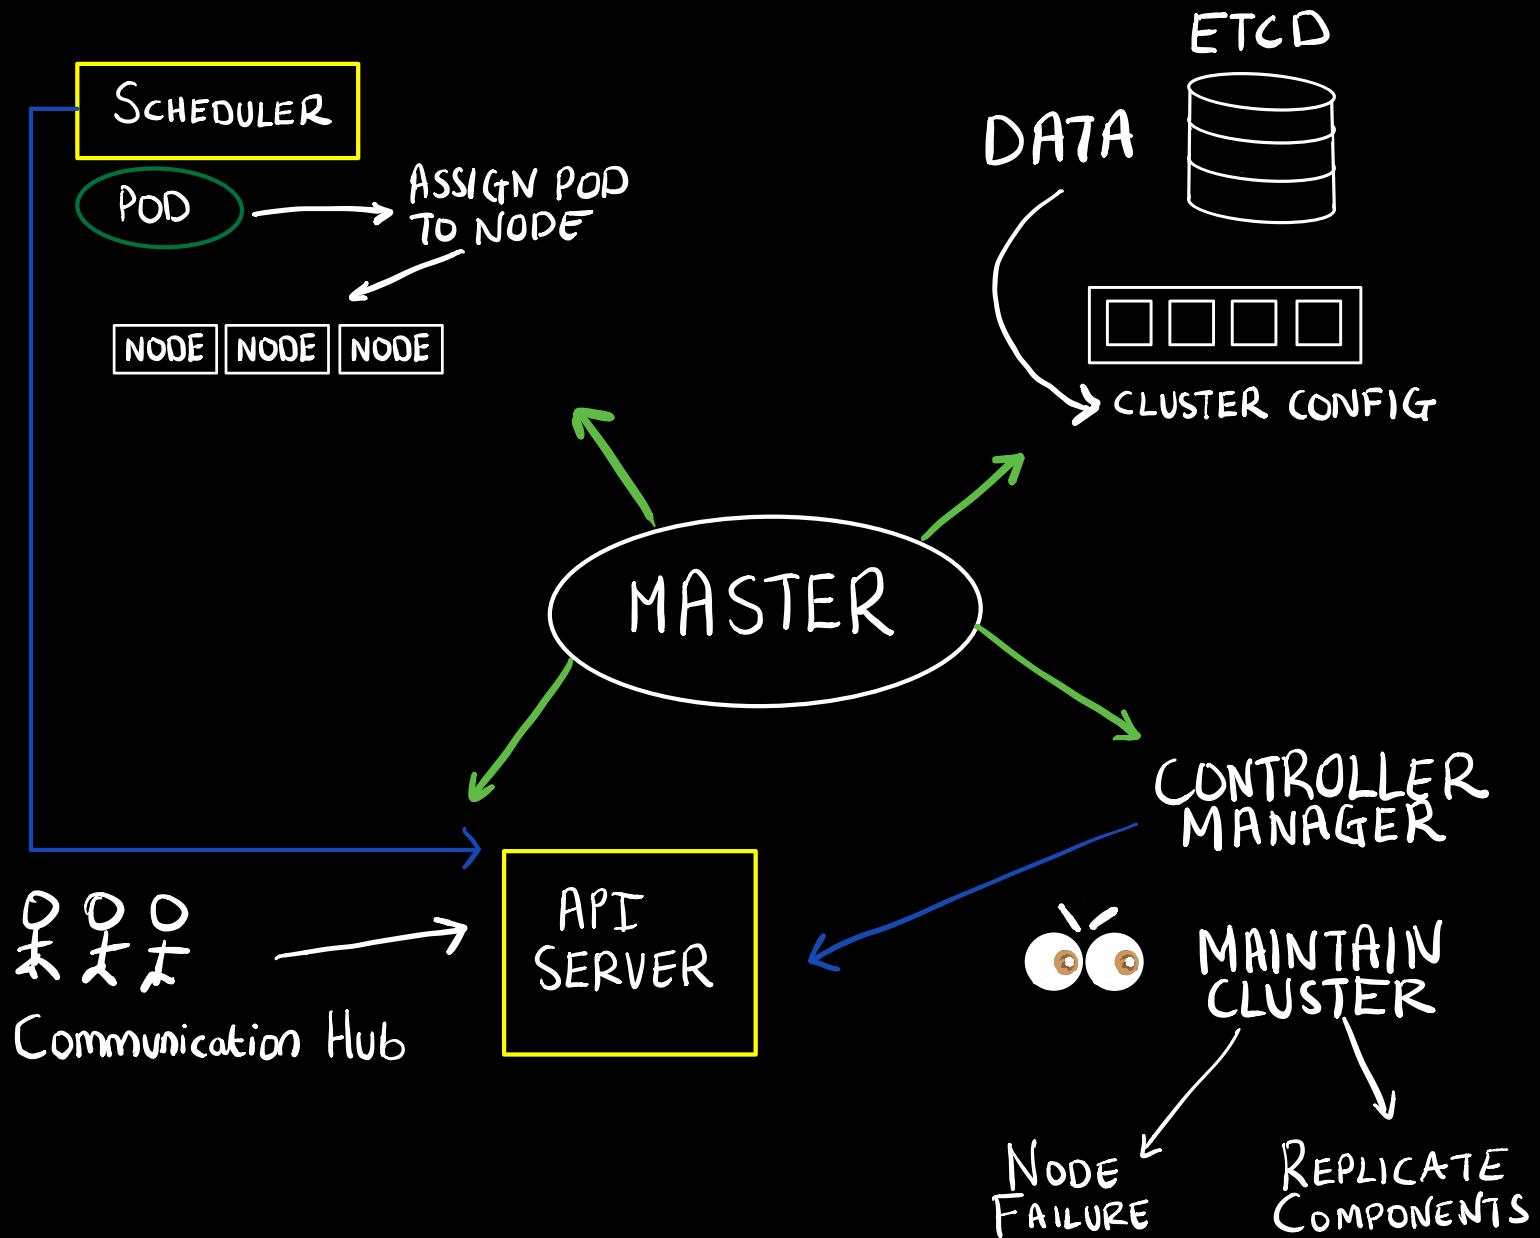


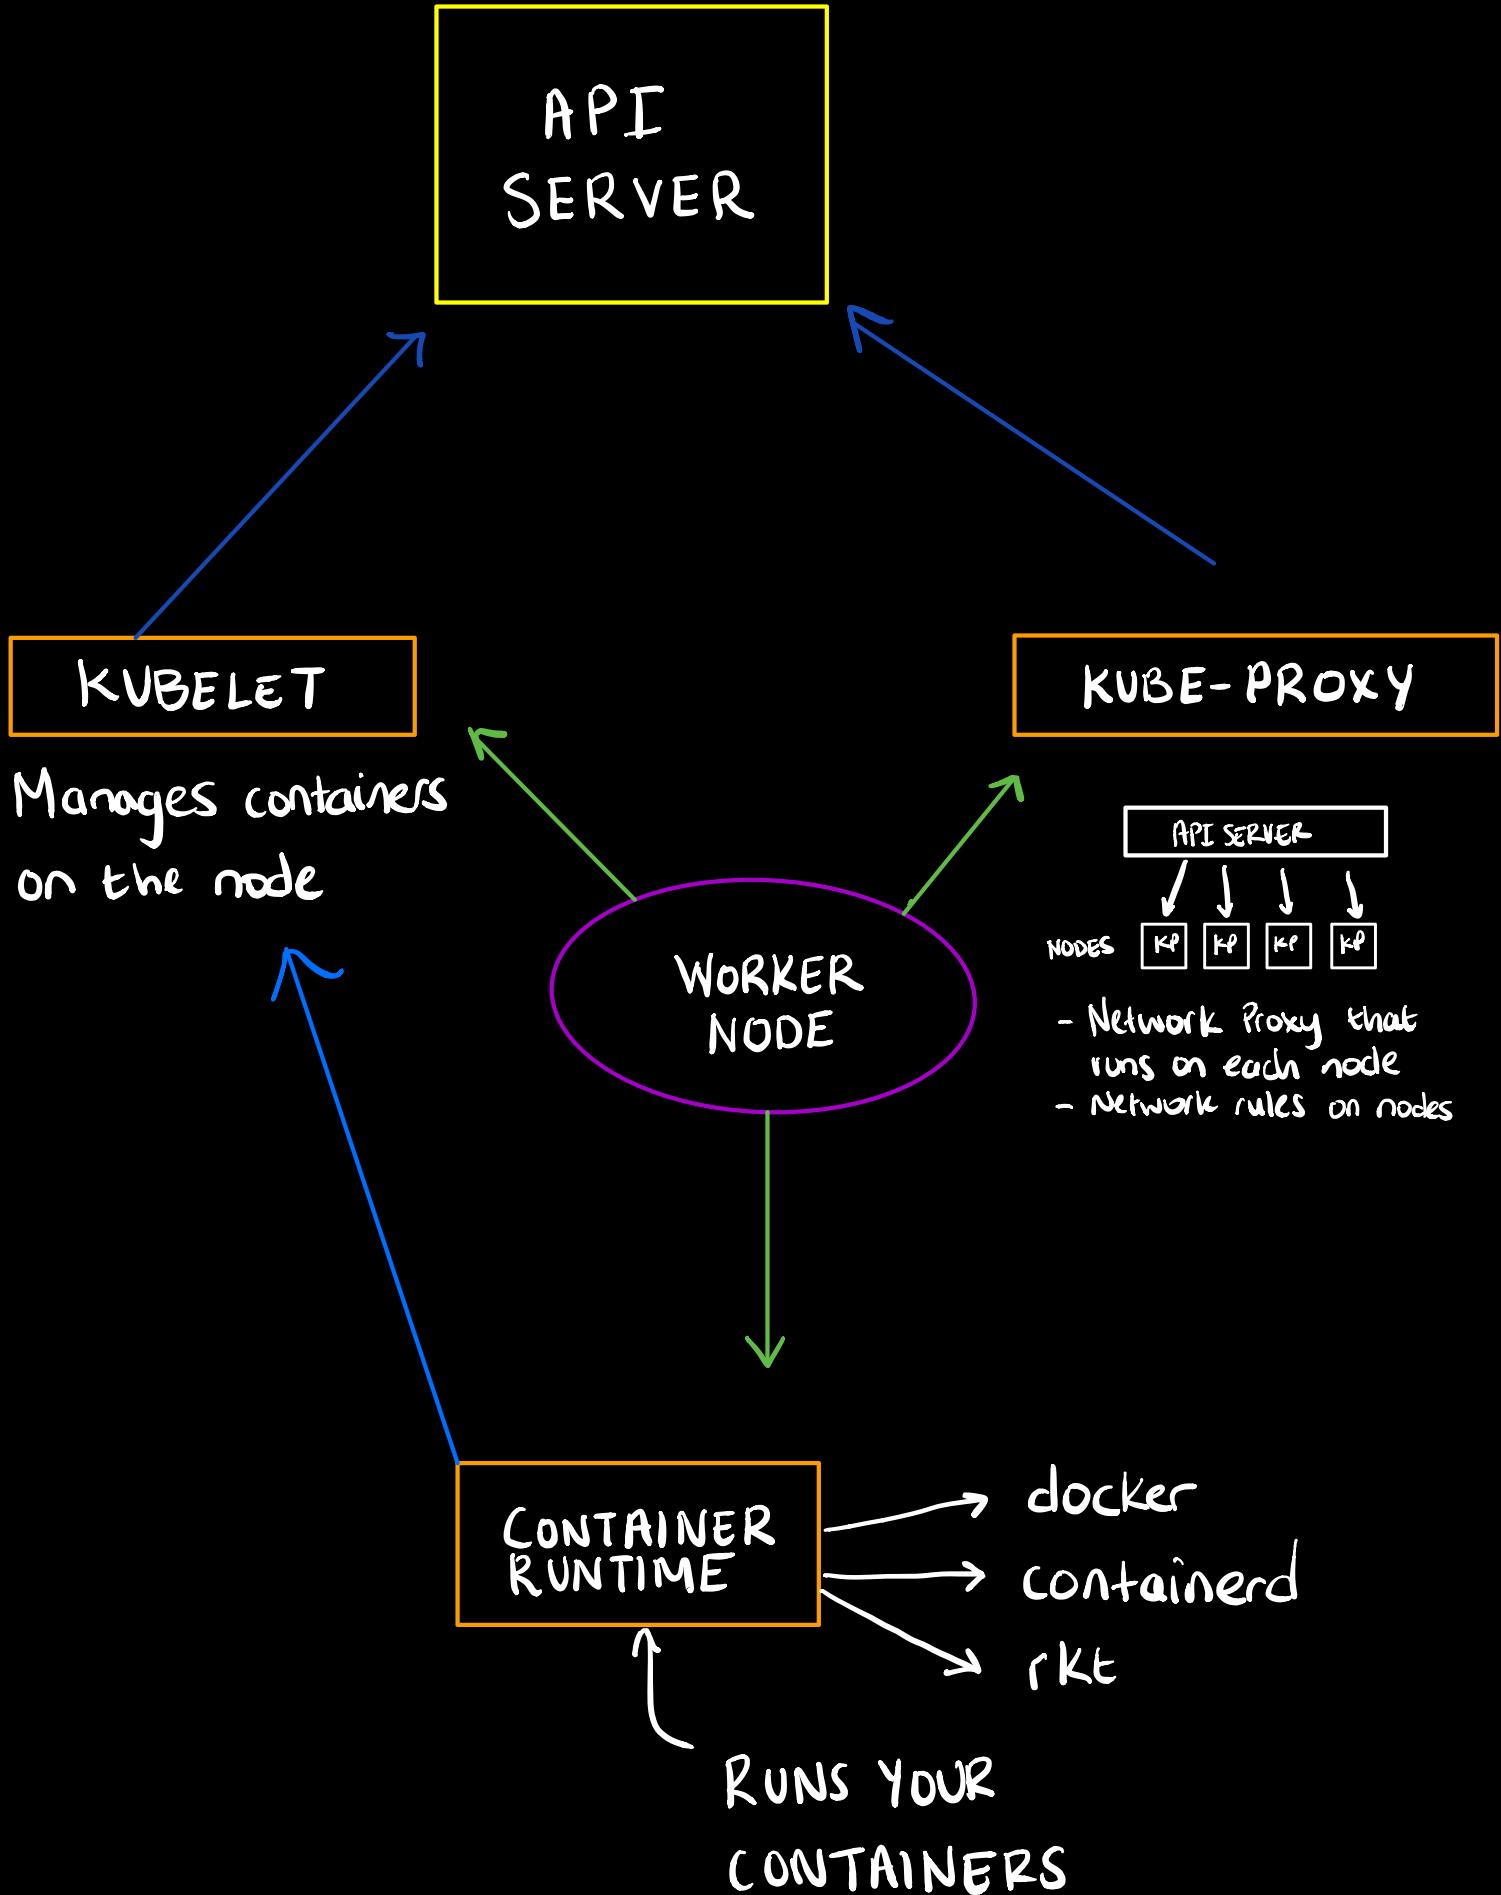
instagram.com/  
**adnans\_techie\_studies**



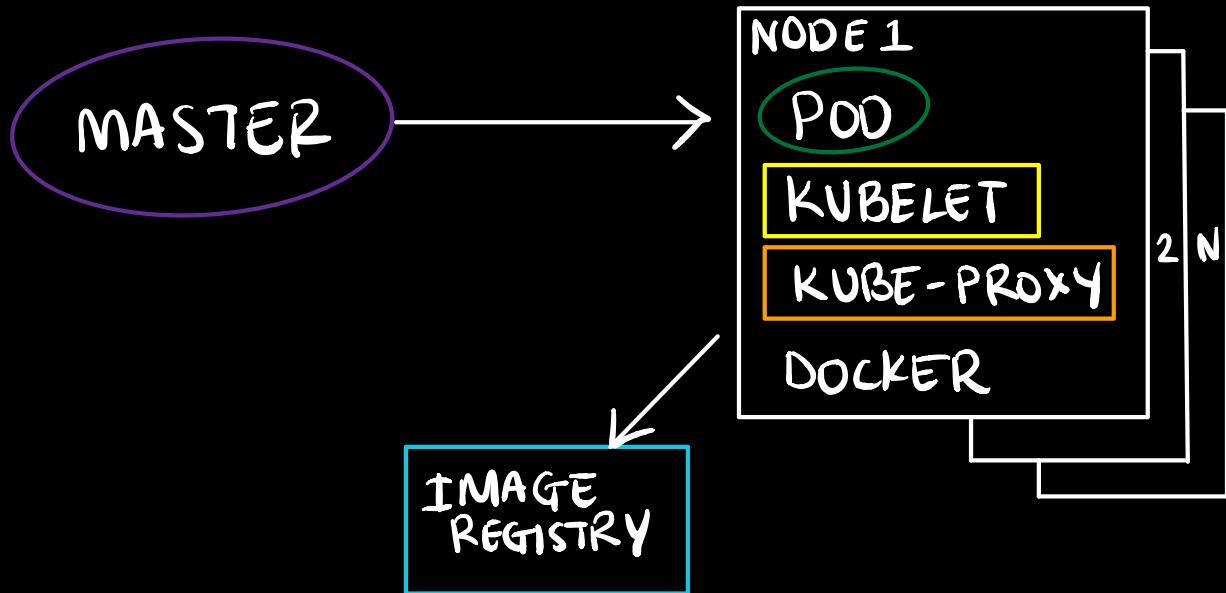
# UNDERSTANDING THE KUBERNETES ARCHITECTURE

# CLUSTER ARCHITECTURE



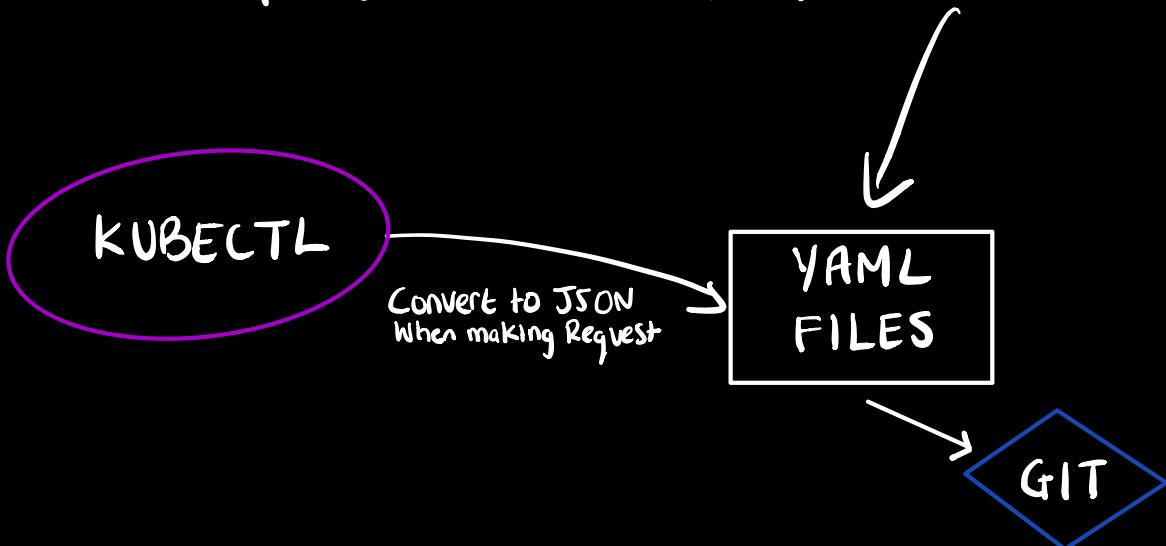


# APPLICATION RUNNING ON K8S



# API PRIMITIVES

API Server is the only one that communicates with etcd  
Every single component speak with API server only not to each other  
Objects like pods and services are declarative intents



# YAML FILE COMPOSITION

**API VERSION**

→ clear consistent view  
of resources

**KIND**

→ The kind of object you want to create

- Pod
- Deployment
- Job

**METADATA**

→ uniquely identify the object

Name  
String

↓  
UID

↓  
namespace

**SPEC**

Container image  
volume exposed ports

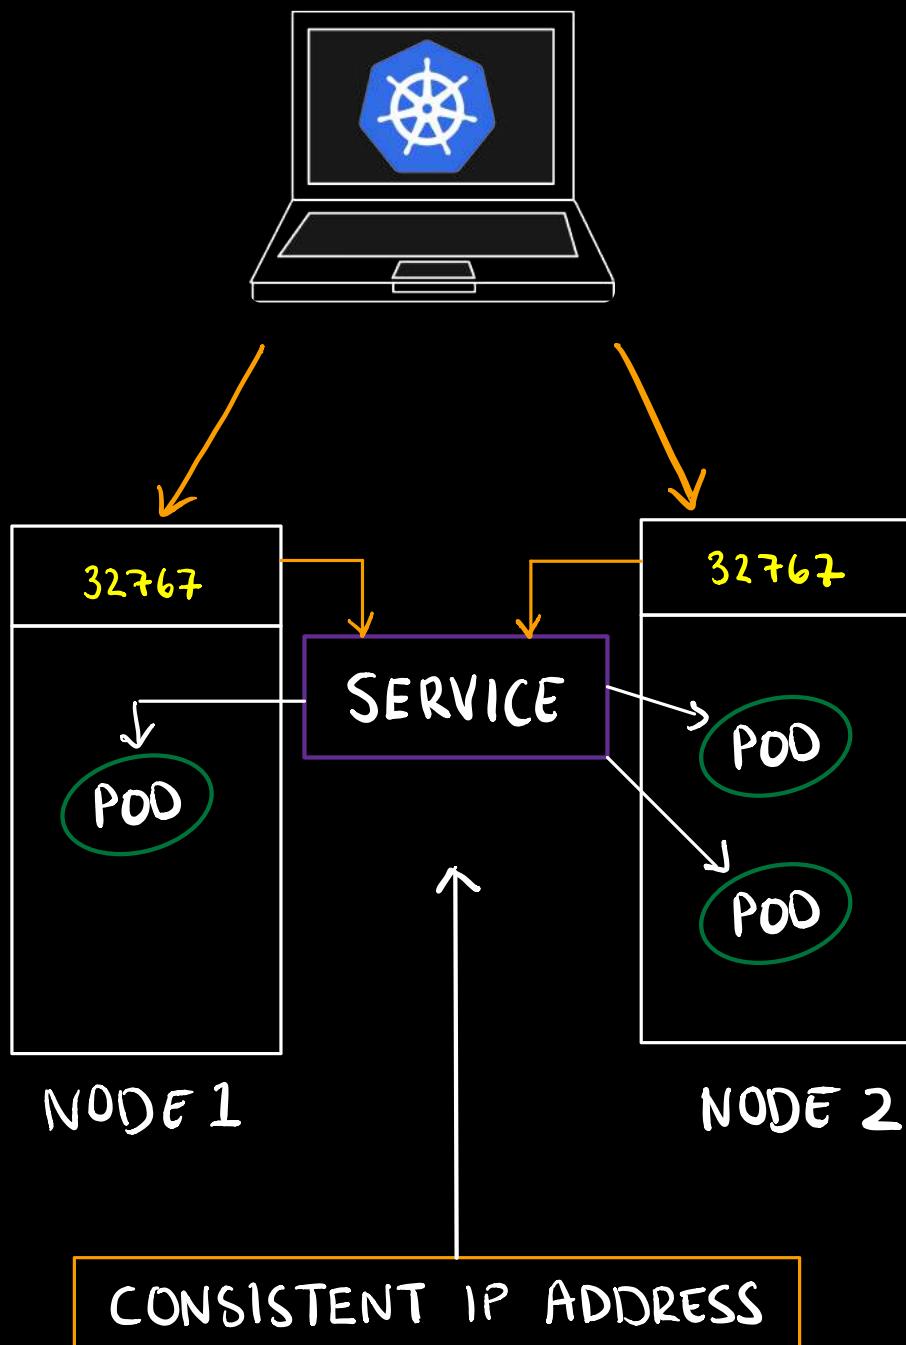
**STATUS**

→ State of the object

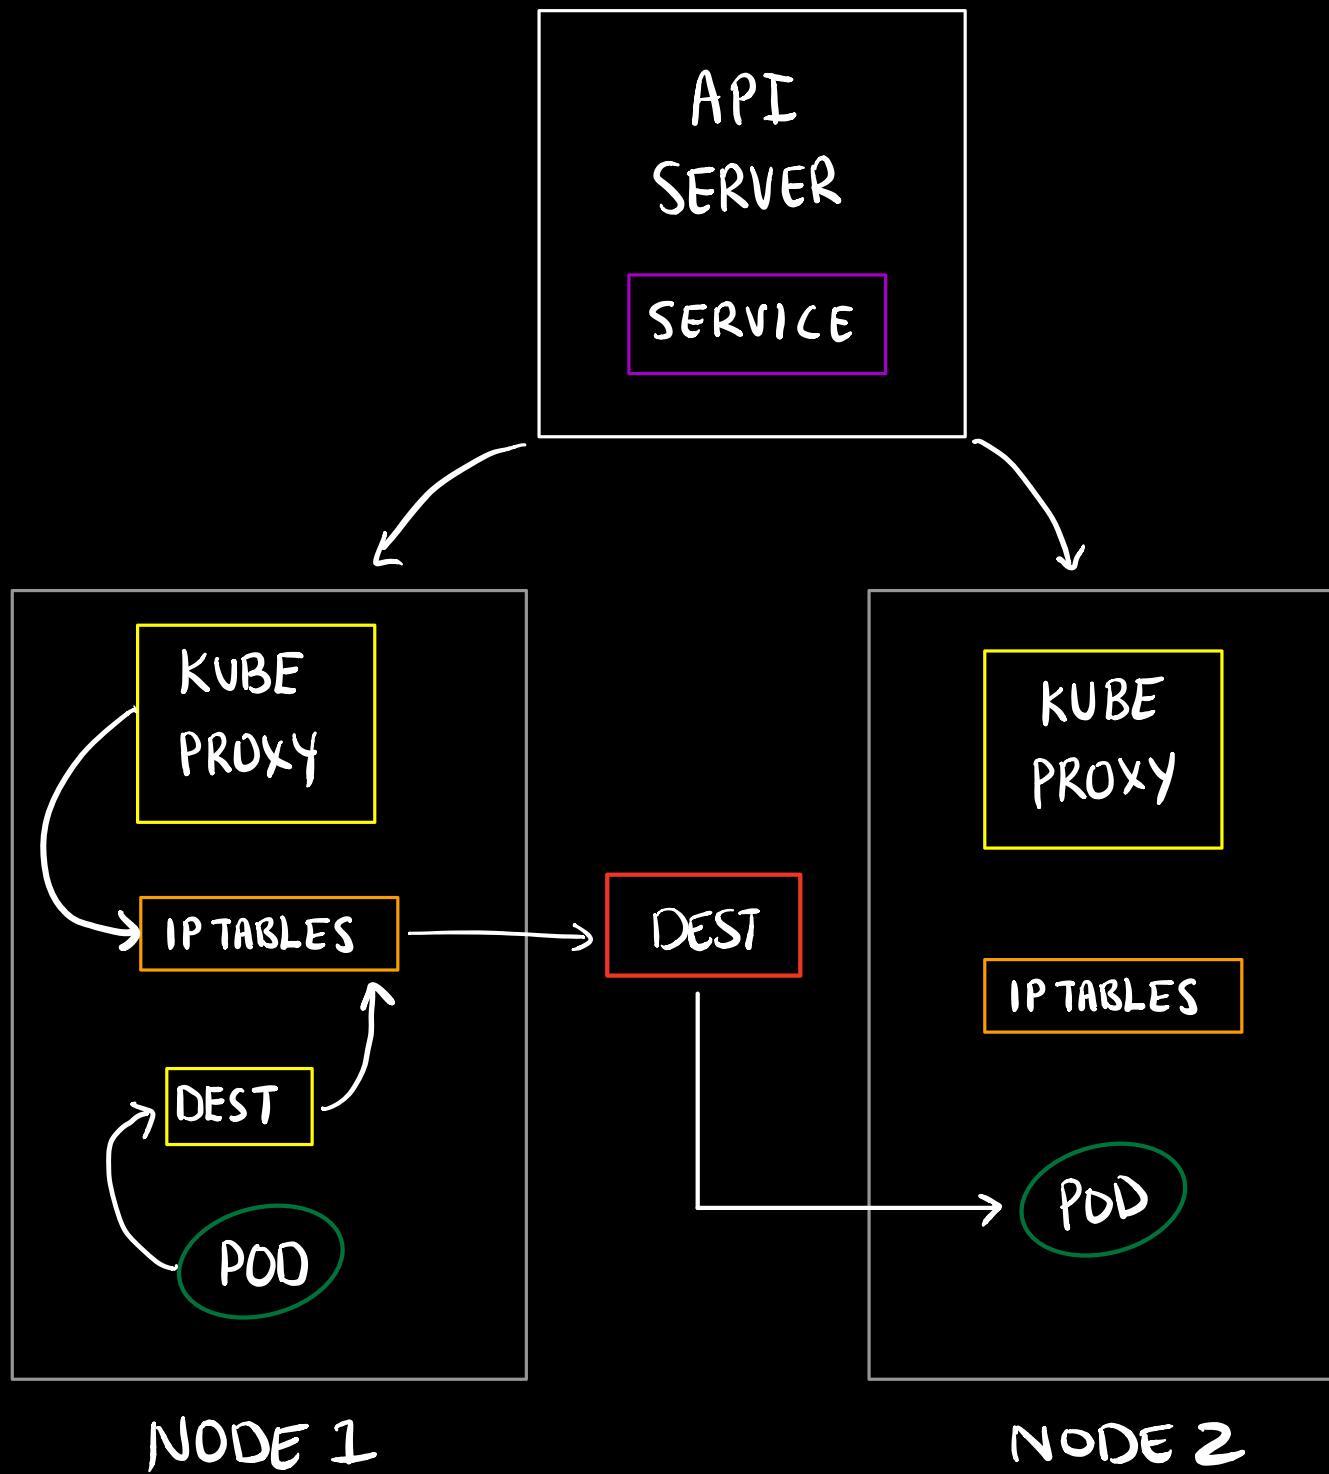
↓  
match desired states

# SERVICES AND NETWORK PRIMITIVES

SERVICES ALLOW YOU TO DYNAMICALLY  
ACCESS A GROUP OF REPLICA PODS



# KUBE-PROXY



KUBE-PROXY HANDLES THE TRAFFIC ASSOCIATED WITH A SERVICE BY CREATING IP TABLE RULES



# BUILDING THE KUBERNETES CLUSTER

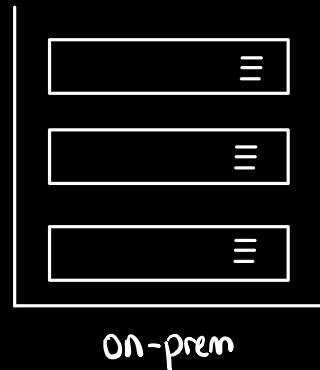
# RELEASE BINARIES, PROVISIONING & TYPES OF CLUSTERS

## Picking the right Solution



CLOUD

VS



ALSO

### CUSTOM

- Install Manually
- Configure your own network fabric
- Locate the release binaries
- Build your own Images
- Secure cluster config

VS

### Pre-Built

- Minikube
- Minishift
- Microk8s
- Ubuntu on LXD
- AWS, Azure, GCP

✓ using minikube locally on MacOS

# INSTALLING KUBE MASTER AND NODES

## MASTER + WORKERS

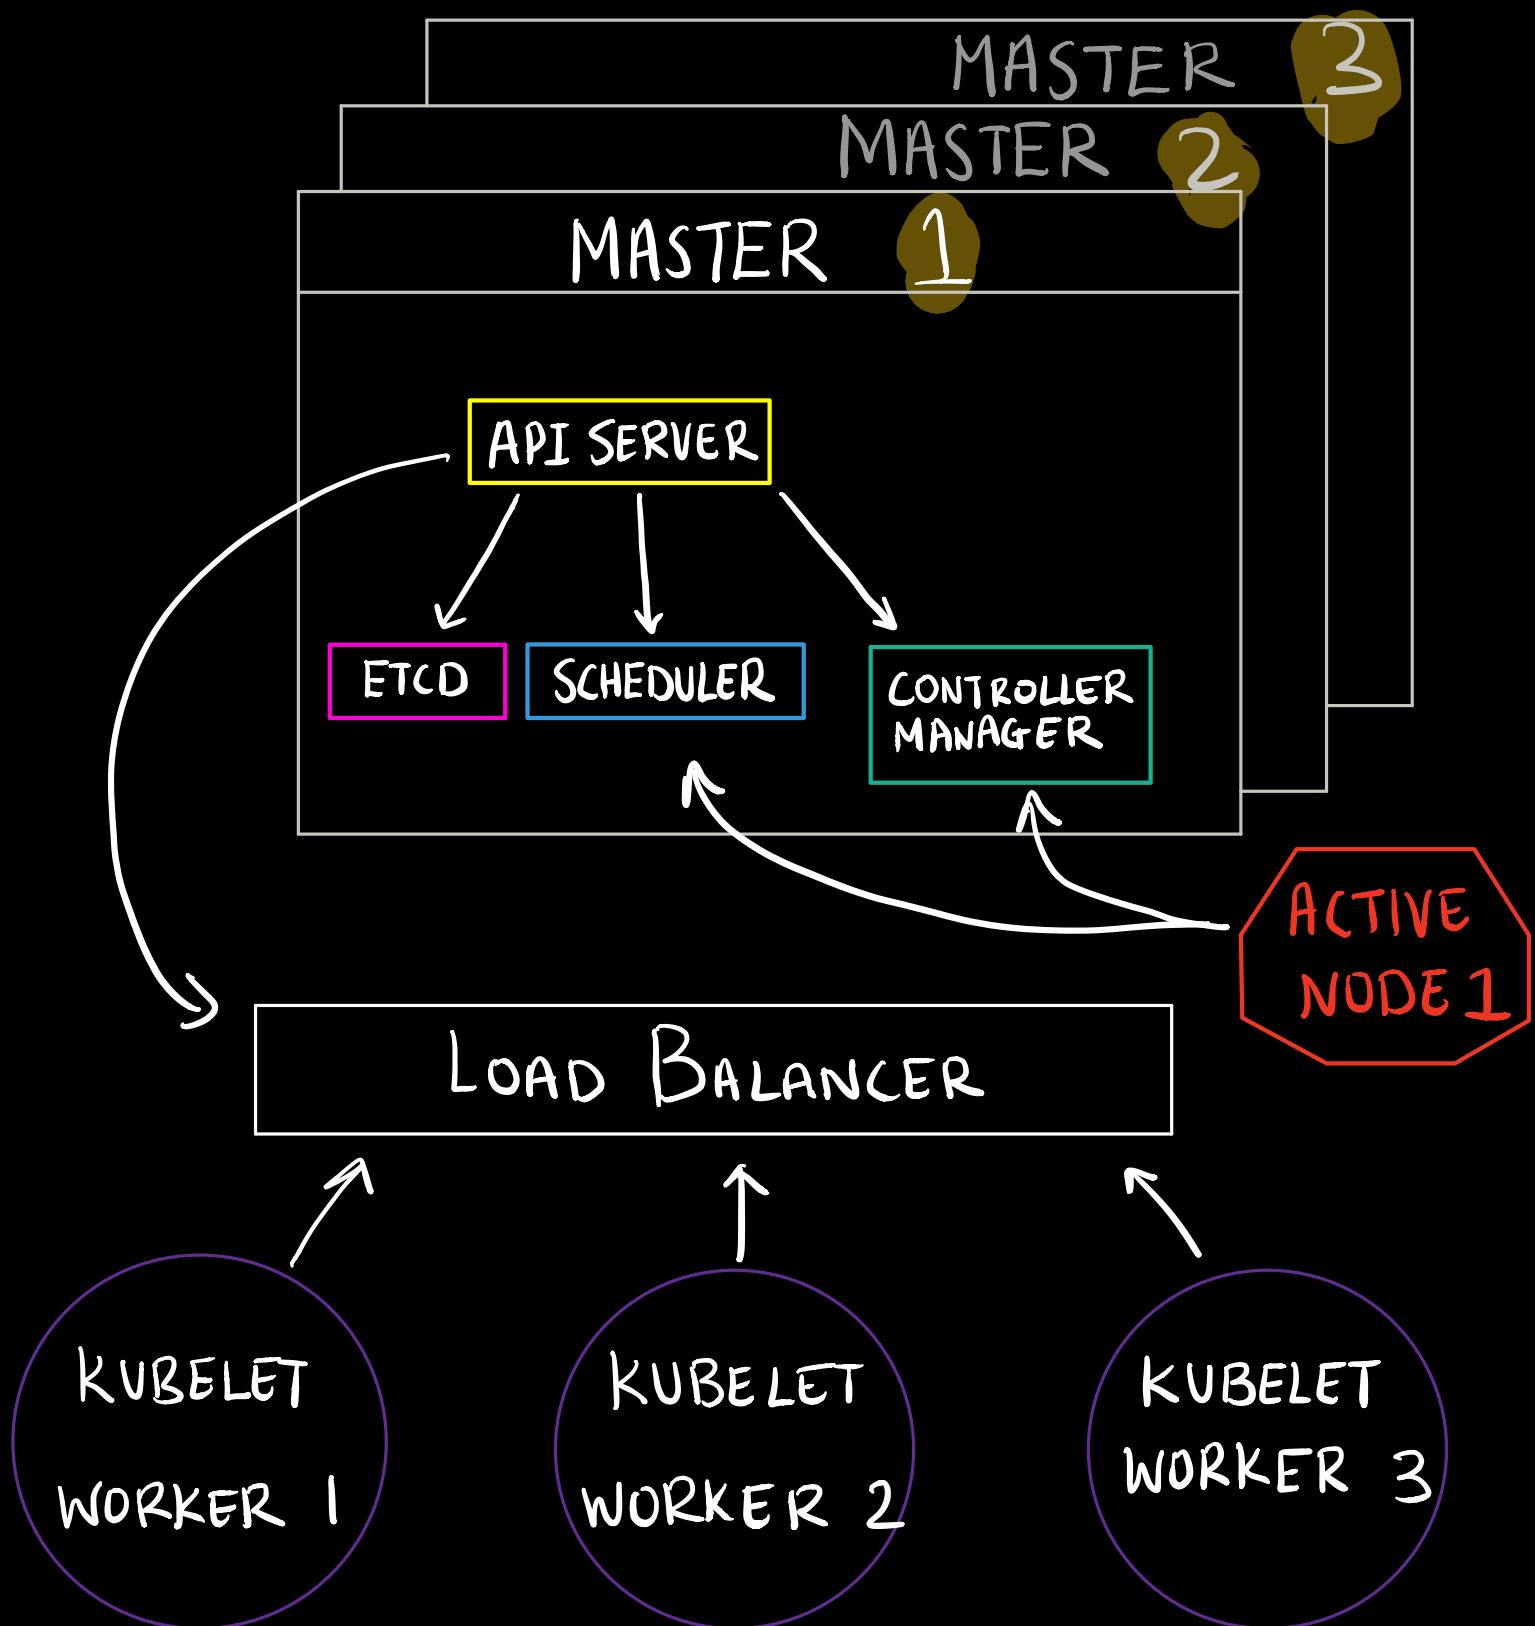
- ① DOCKER + KUBE → GPG KEY  
→ ADD REPOS
- ② UPDATE PACKAGES
- ③ INSTALL Docker, kubelet, kubeadm, kubectl
- ④ Modify bridge adapter settings

## MASTER ONLY

- ① Initialise Cluster
- ② Make directory for kss
- ③ Copy kube config
- ④ Change ownership of config
- ⑤ Apply Flannel CNI

# BUILDING HIGHLY AVAILABLE CLUSTER

ALL components can be replicated, but only certain can operate simultaneously



The controller manager and the scheduler actively watch the cluster state and take action when it changes

SCHEDULER



CM



CLUSTER

ARE WE IN

CHARGE?

SCHEDULER



CM



NO WE ARE!

→ Leading to duplicate resources or corruption.

HOW DO WE  
DECIDE?

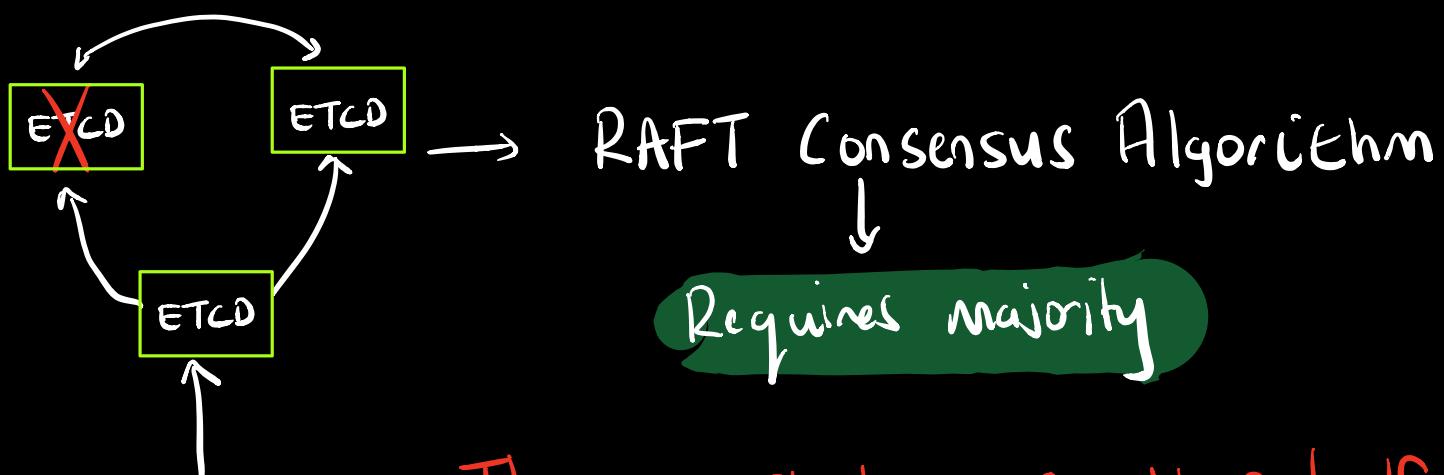
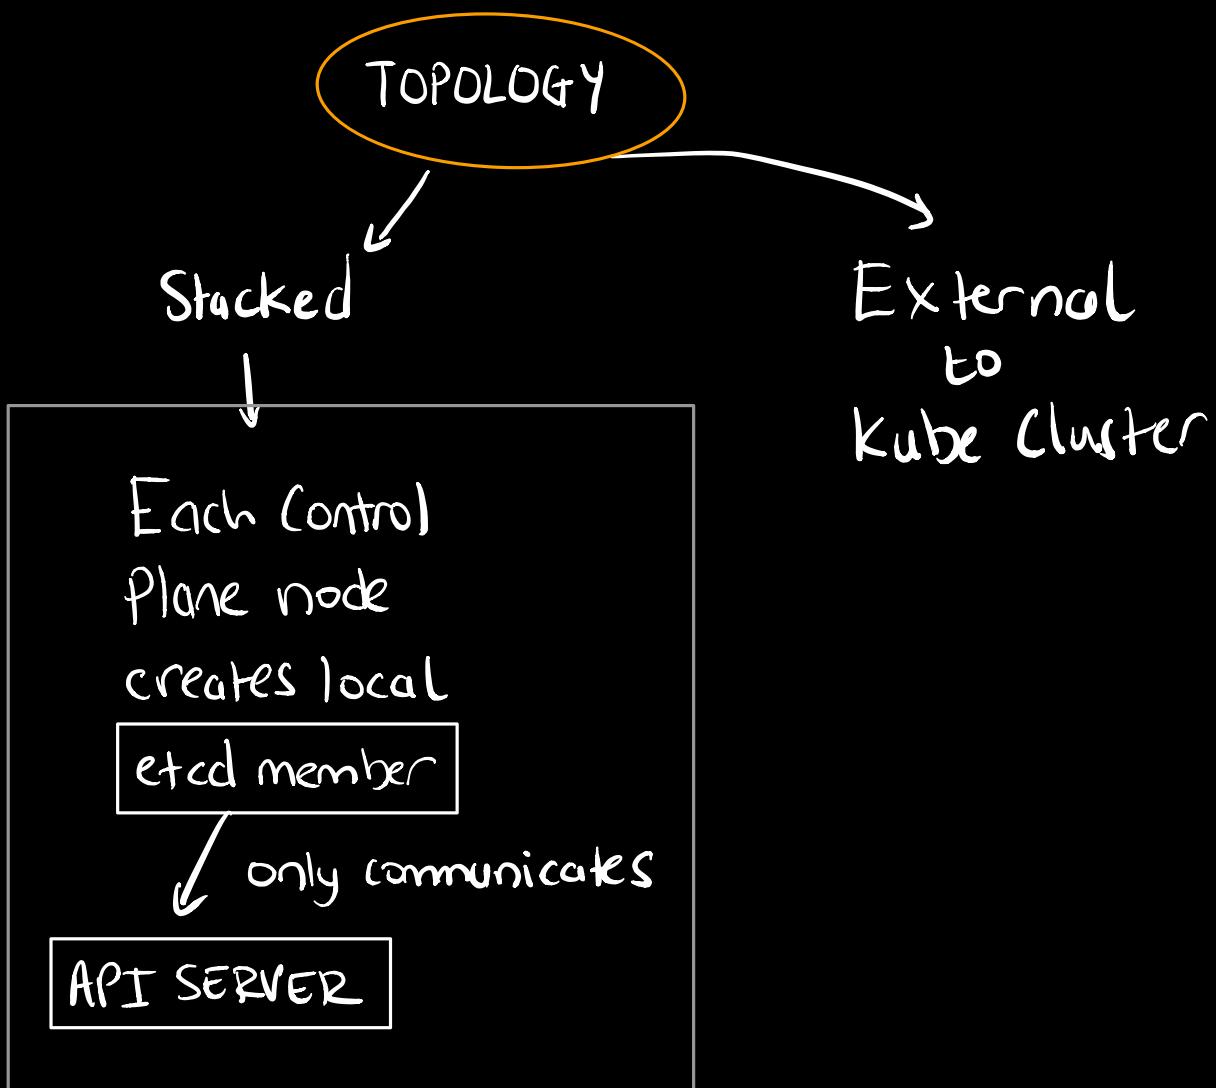
LEADER  
ELECT  
OPTION

→ creates endpoint resource

↑ see in scheduler YAML

↑ holderIdentity

# REPLICATING ETCD

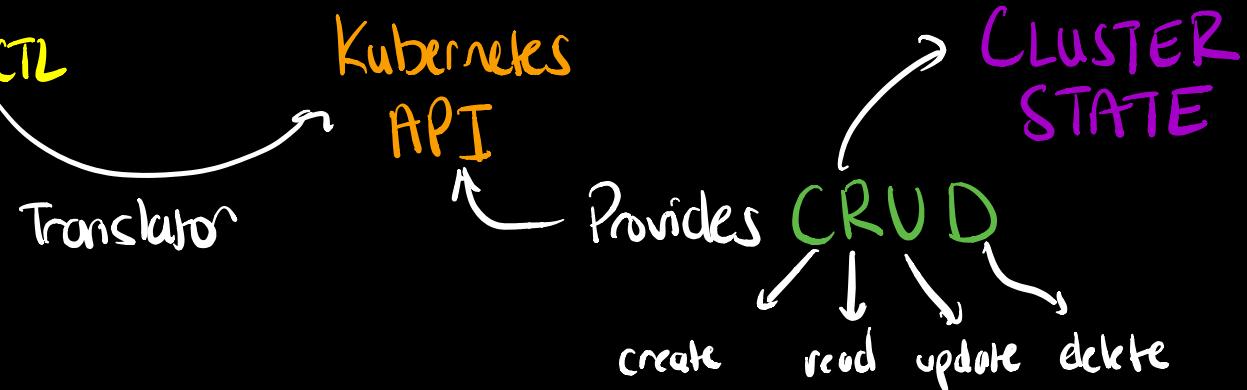


There must be more than half taking place in the state change  
∴ there must be odd number of nodes

# CONFIGURING SECURE CLUSTER COMMUNICATIONS

→ All communication via HTTPS

→ Kube CTL



Return Response

→ Kube CTL  
CREATE A POD

HTTP POST  
API SERVER

State Stored in ETCD

MULTIPLE PLUGINS

Authentication

Authorisation

Admission

Validation

calls to determine request

Can this User perform this action?

READ → Skip

create  
Modify  
delete

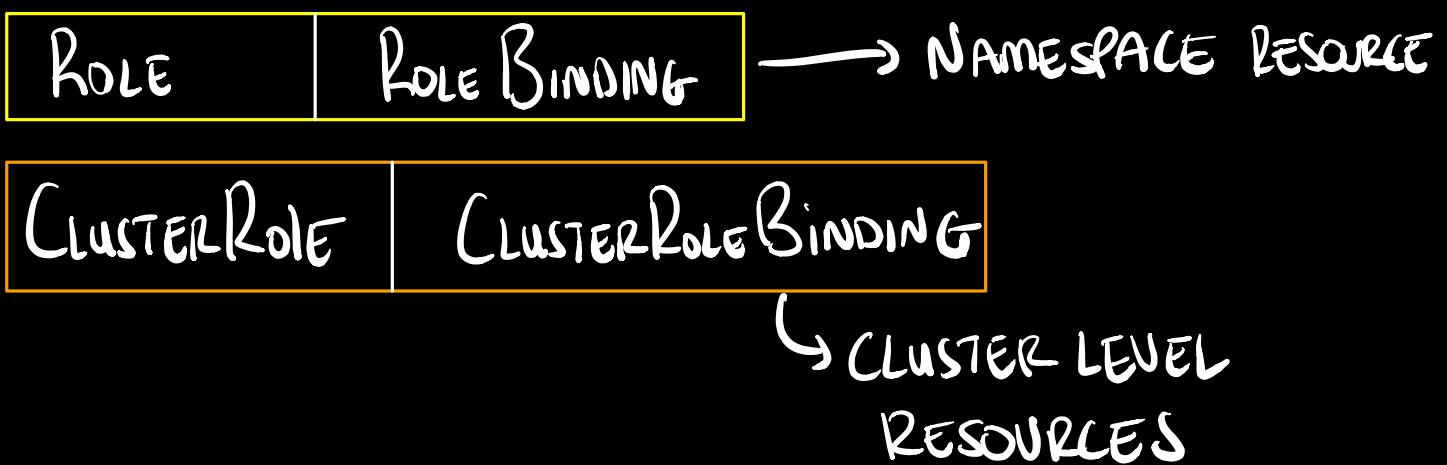
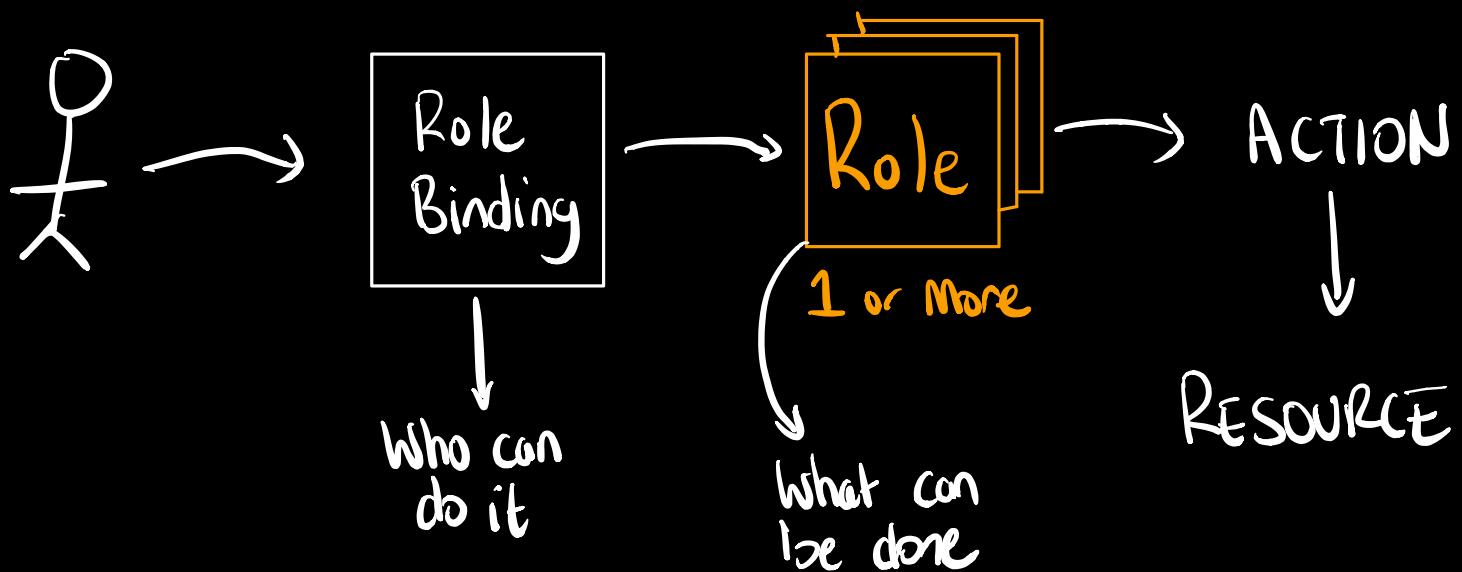
HTTP

CERTIFICATE

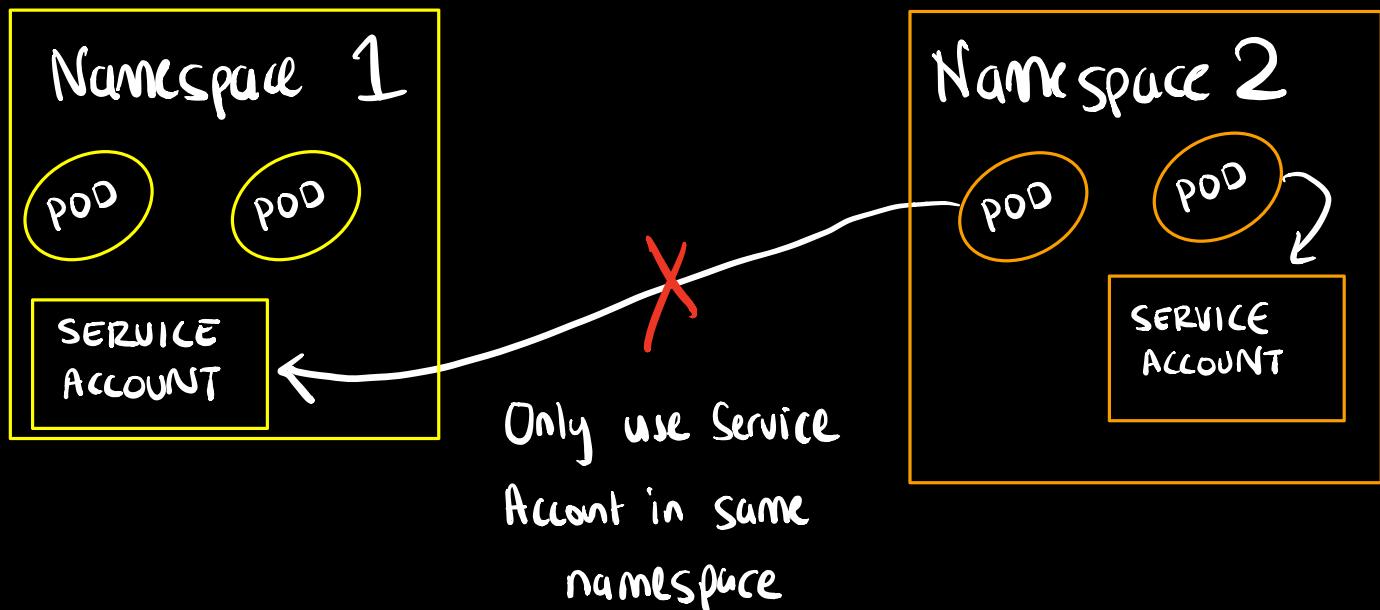
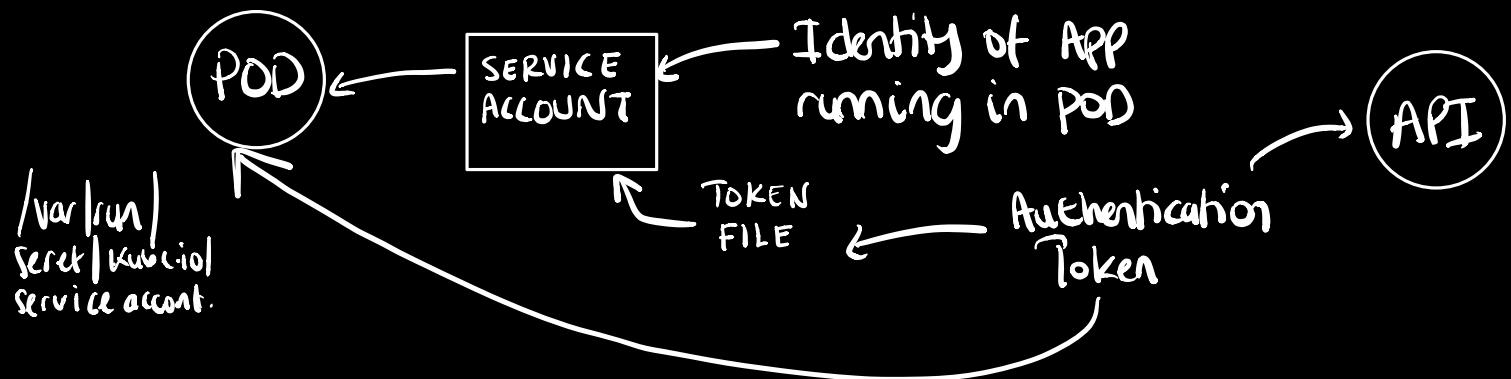
HEADER

# ROLES AND ACCESS

RBAC is used to prevent unauthorised users from modifying the cluster state



# SERVICE ACCOUNT



# RUNNING END TO END TESTS ON CLUSTER

Performance  
and  
Response of  
Application

Poor Cluster  
Performance

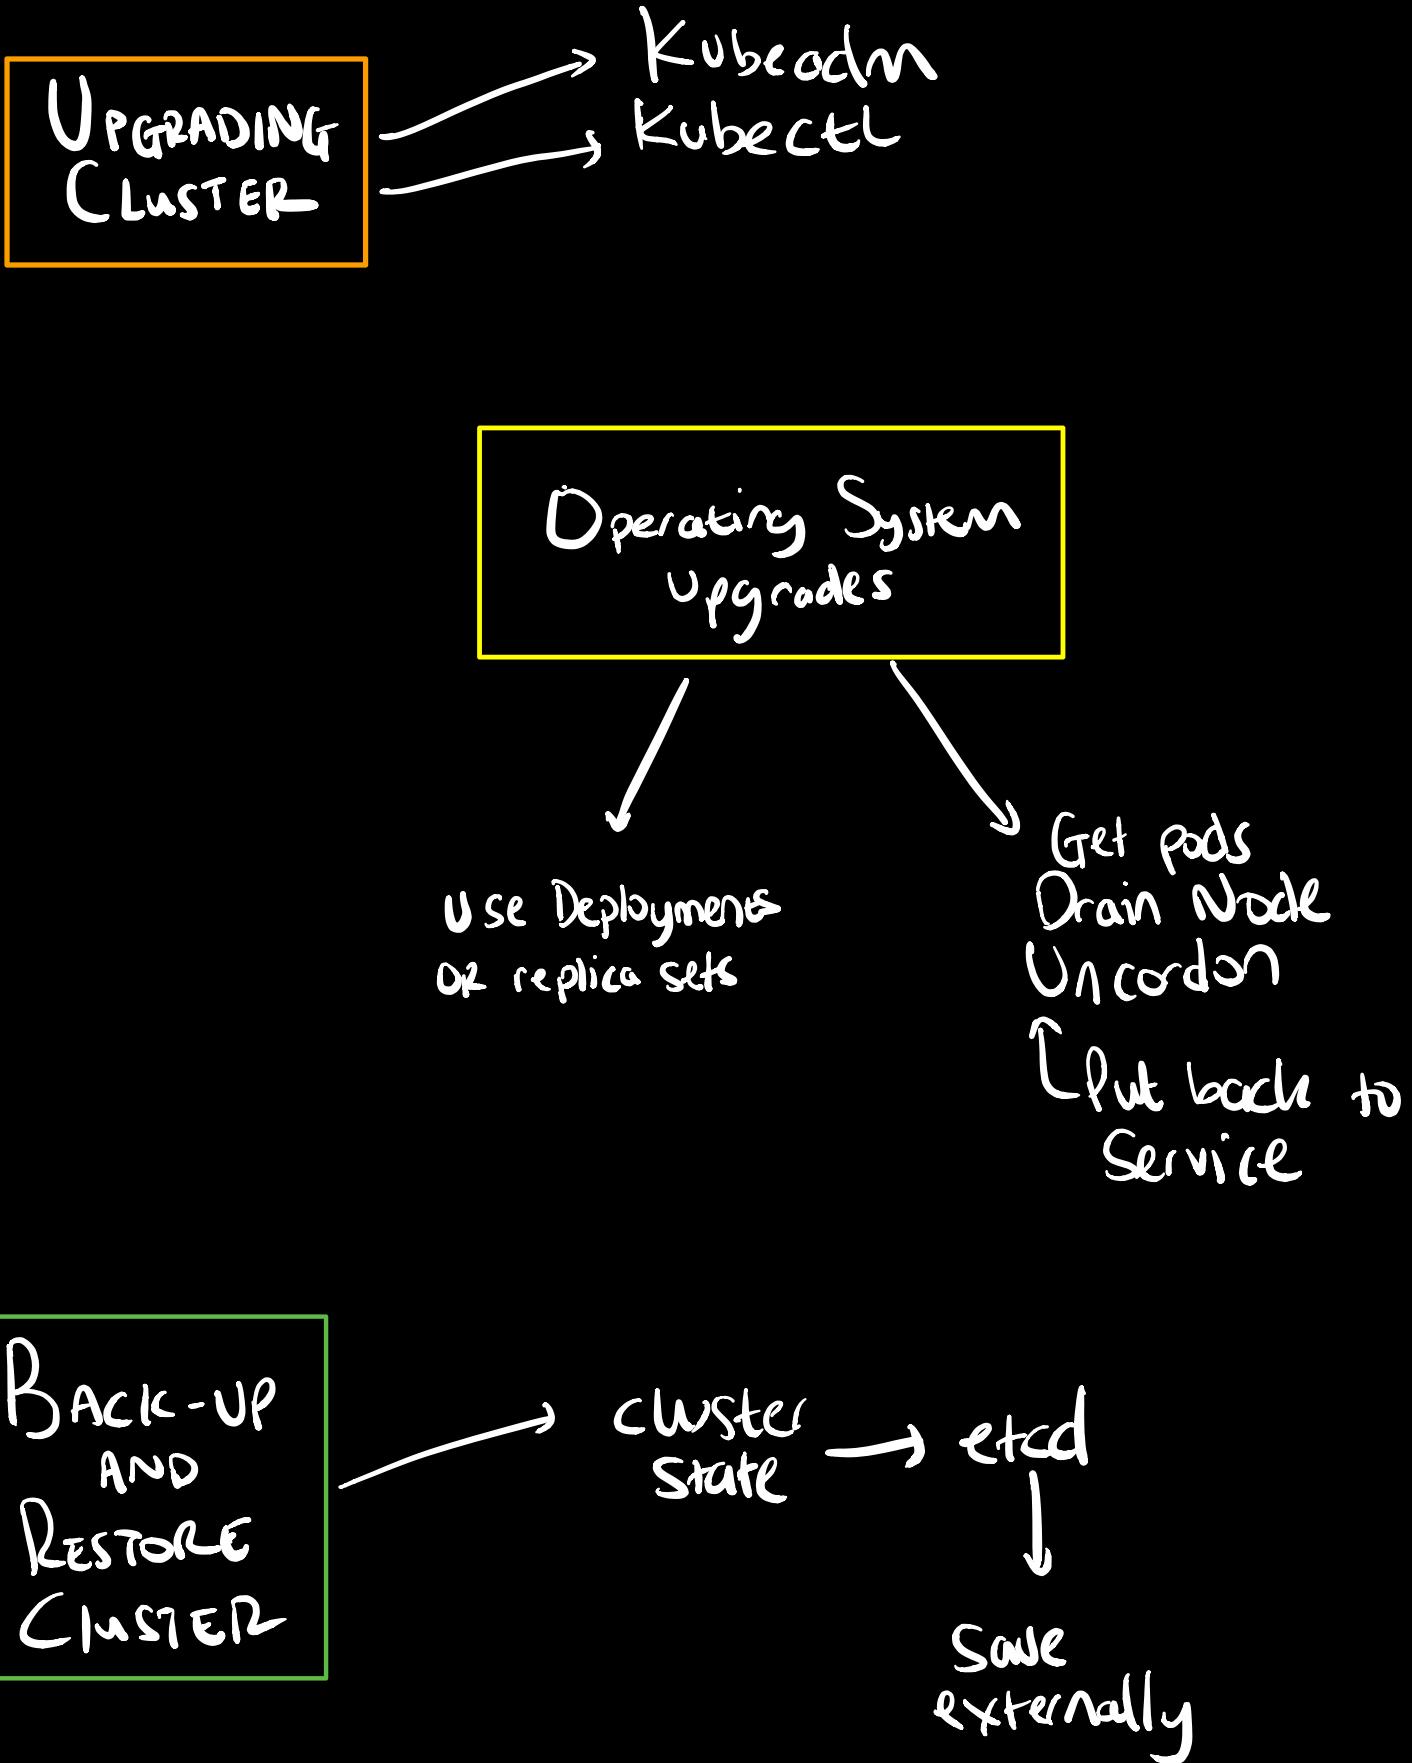


kubeTest

Example  
Tests

- ✓ Deployments can run
- ✓ Pods can run
- ✓ Pods can be directly accessed
- ✓ Logs can be collected
- ✓ Commands run from pod
- ✓ Services can provide access
- ✓ Nodes are healthy
- ✓ Pods are healthy

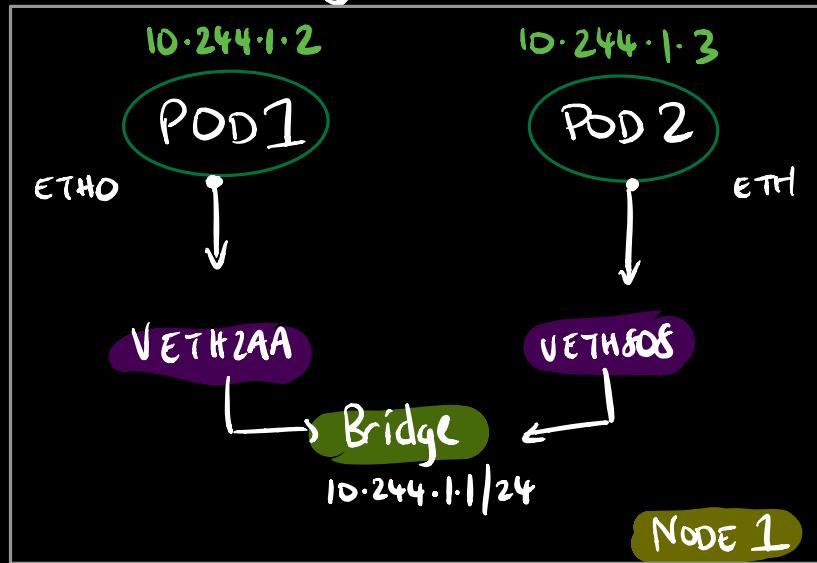
# MANAGING CLUSTER



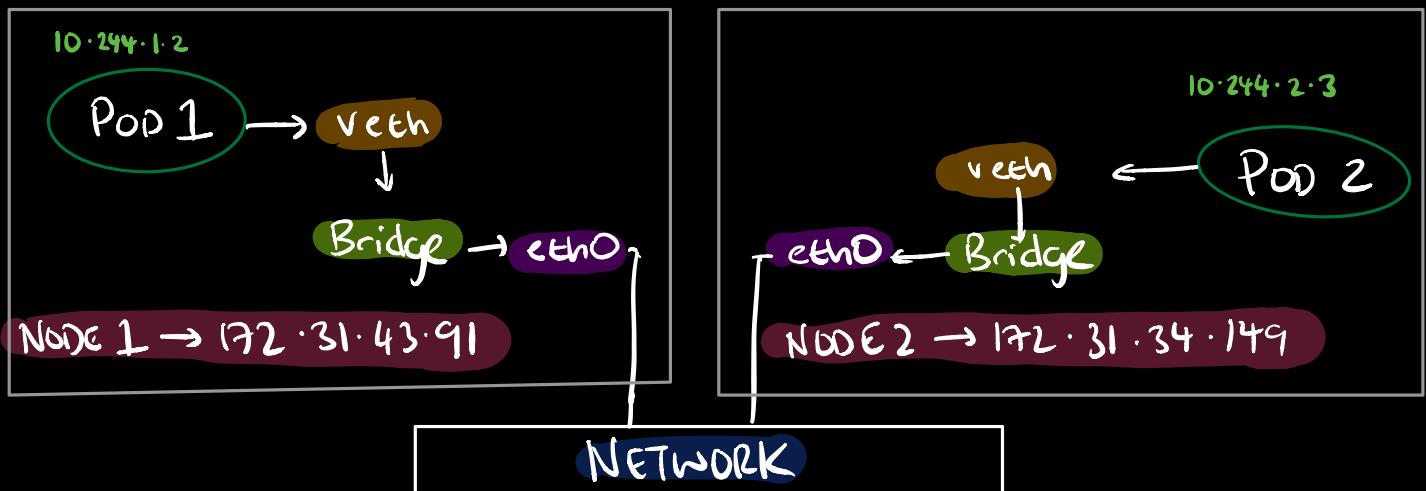
NETWORK  
CLUSTER  
COMMUNICATIONS

# POD AND NODE NETWORKING

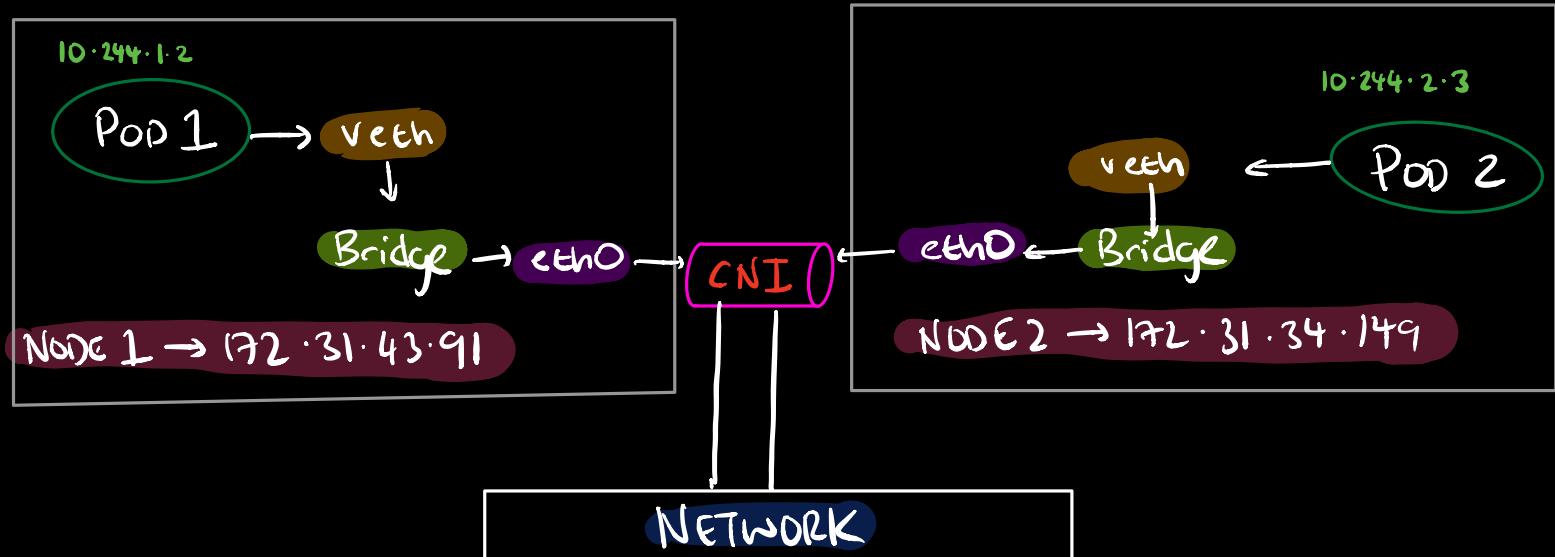
## Networking within a node



## Networking Outside of the Node



# CONTAINER NETWORK INTERFACE



## CNI is a network Overlay

- ↳ Allows building tunnel between nodes
  - ↳ Sits on top of existing networks
    - ↳ Encapsulates Packet
      - ↳ changes SRC & DST



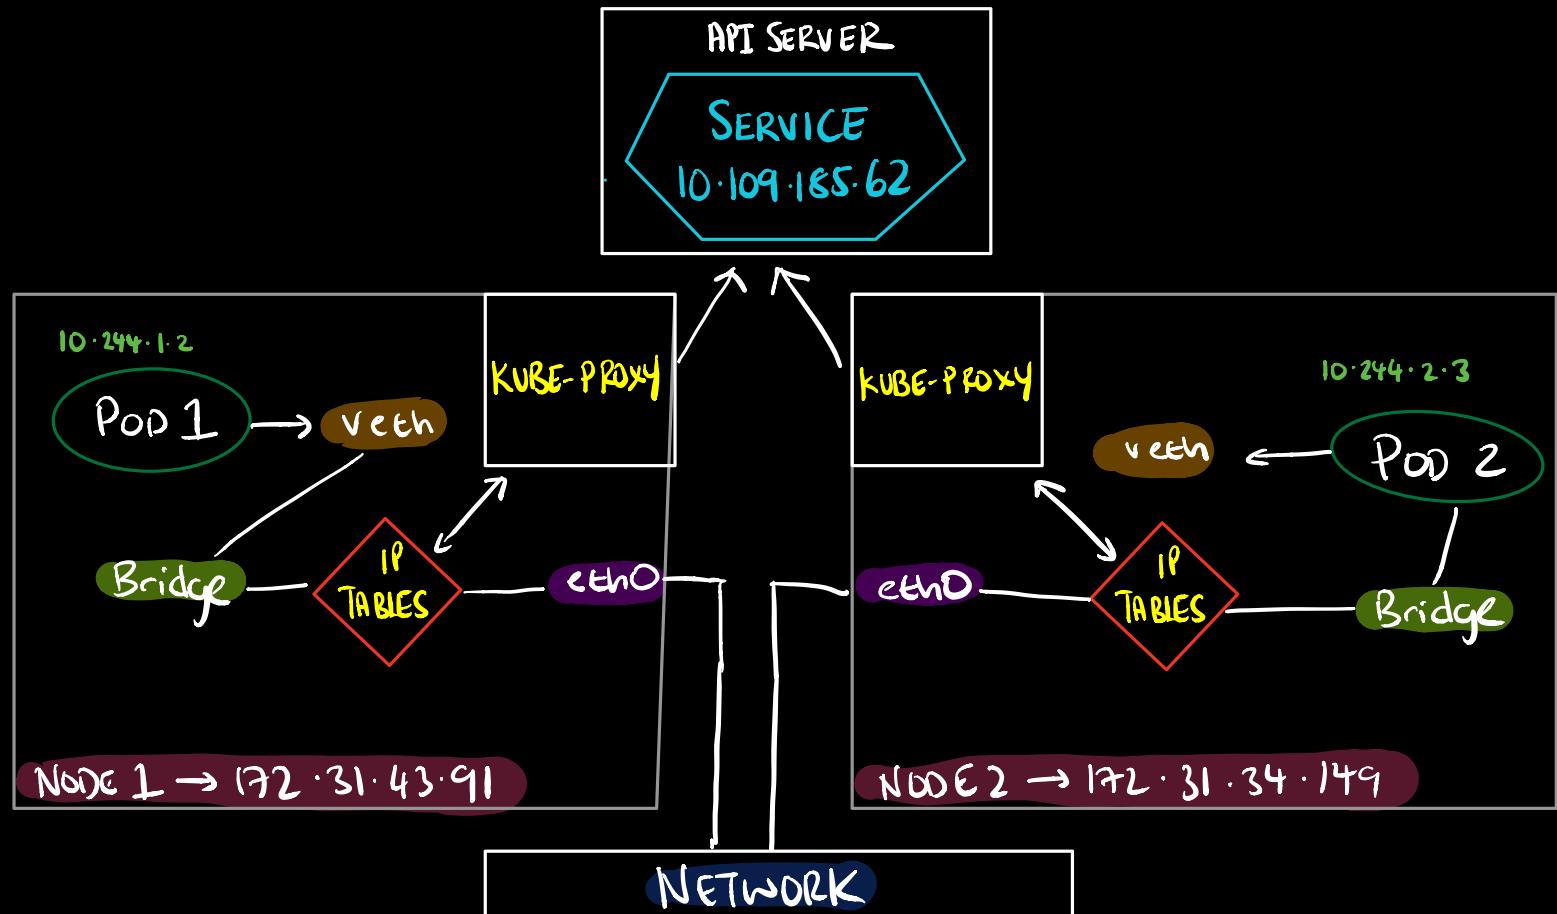
How  
CNI  
DO  
THIS?

- There is a mapping associated in user space
  - ↳ Program all pod ip address's to node IP, when reach other node, de-encapsulate packet and give to bridge

Example CNI  
Calico Flannel

Appears Local To Node!

# SERVICE NETWORKING



→ Pods come and go.

↳ How does cluster keep track?

↳ Services!

↳ Provides virtual interface → Auto assigned to pods behind interface.

Example

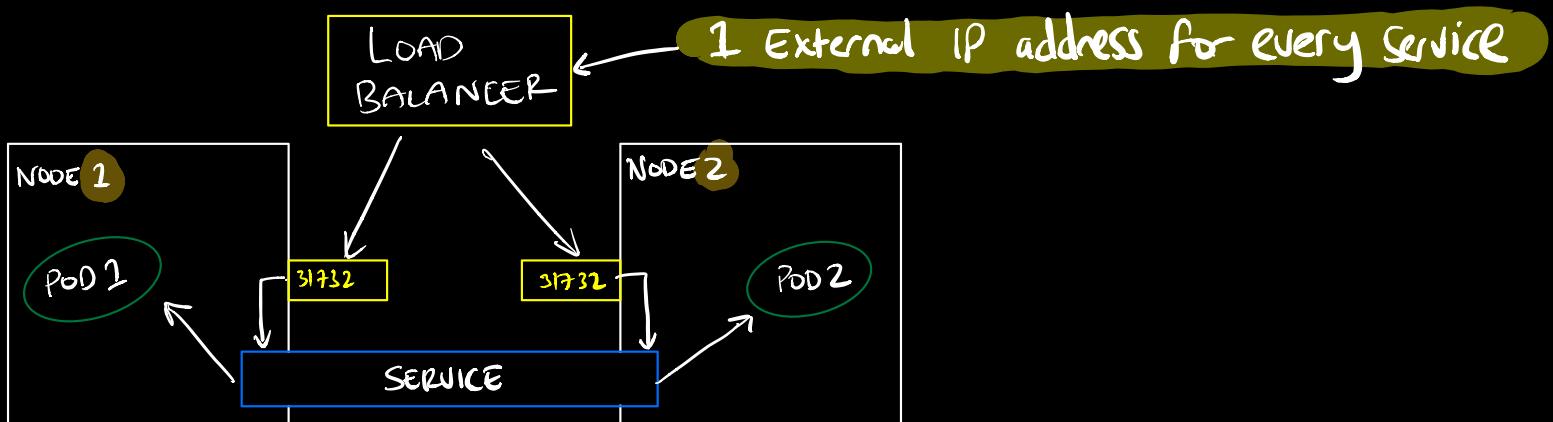
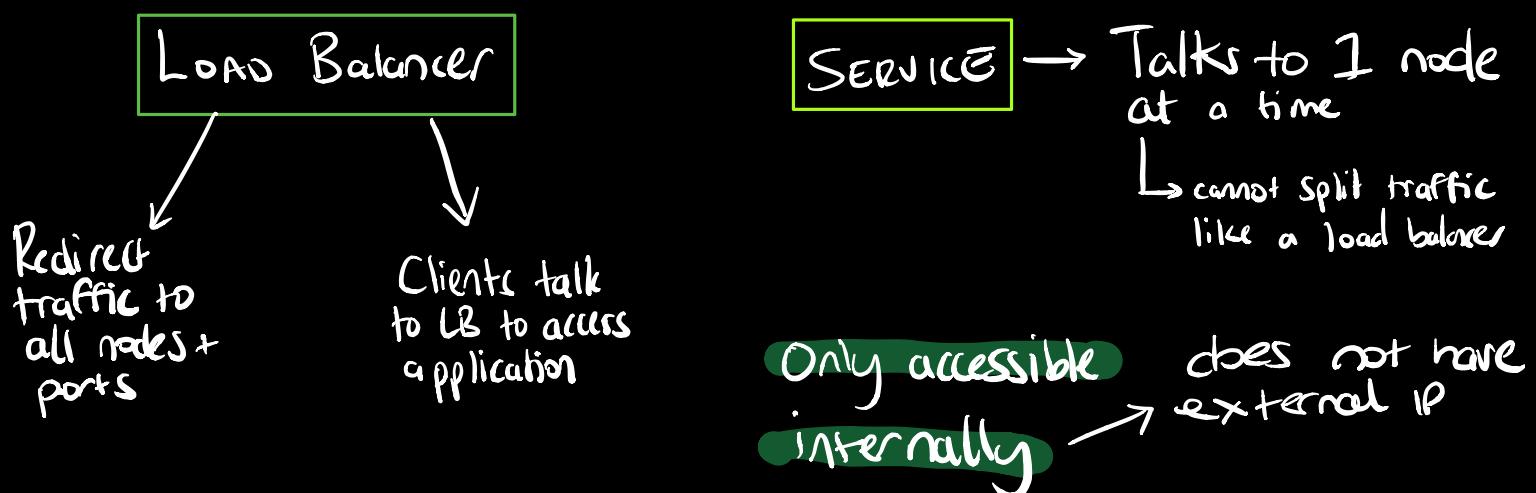
Cluster IP Services

→ Auto created on cluster creation

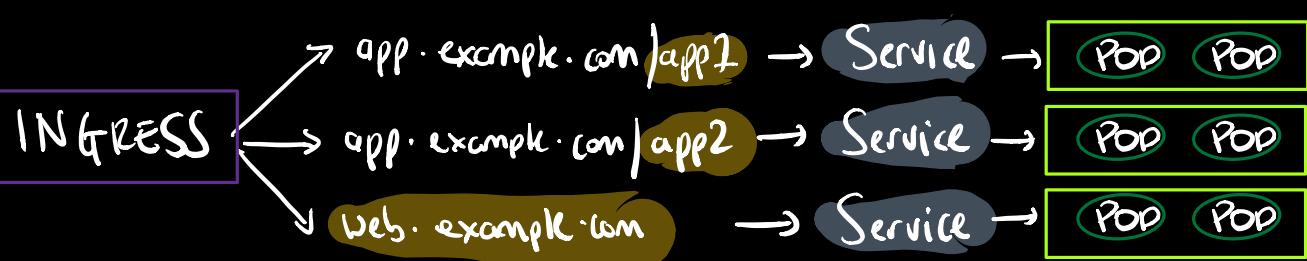
→ Takes care of internal routing

→ No matter where moves other pods know how to communicate to it

# INGRESS RULES AND LOAD BALANCERS



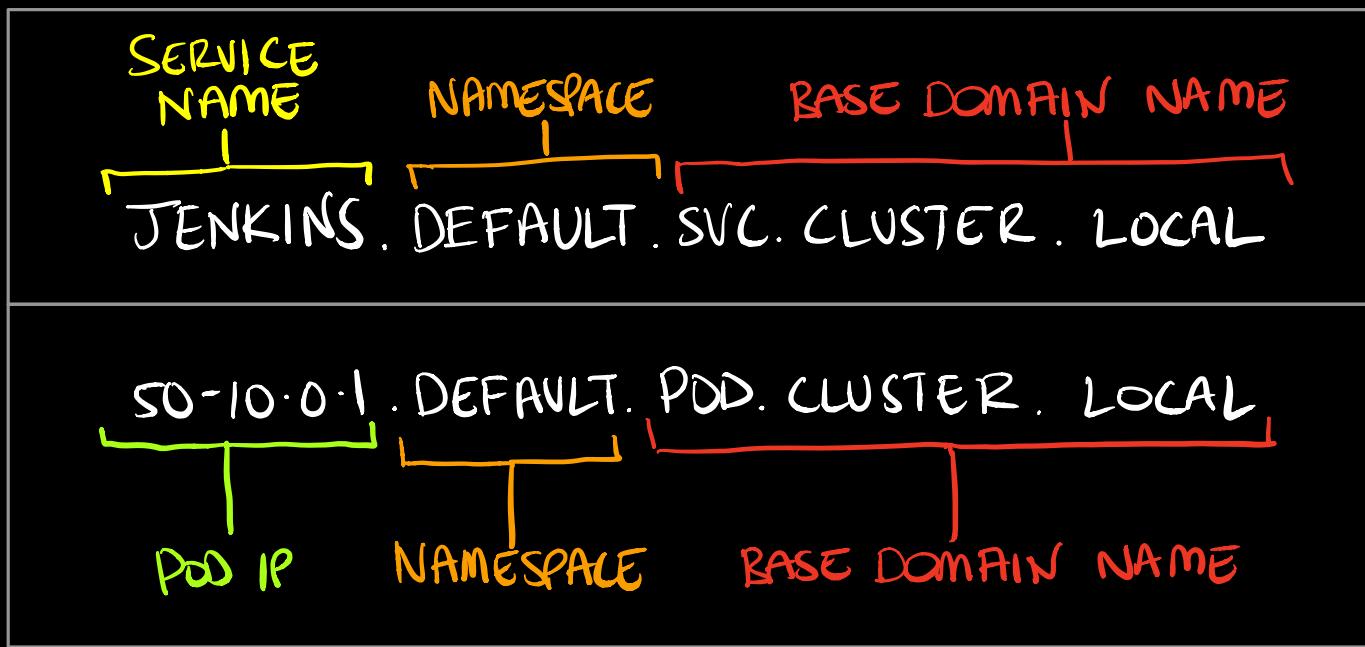
## INGRESS



Access multiple Services with Single IP Address

# CLUSTER DNS

EVERY SERVICE DEFINED IN THE CLUSTER IS ASSIGNED A DNS NAME

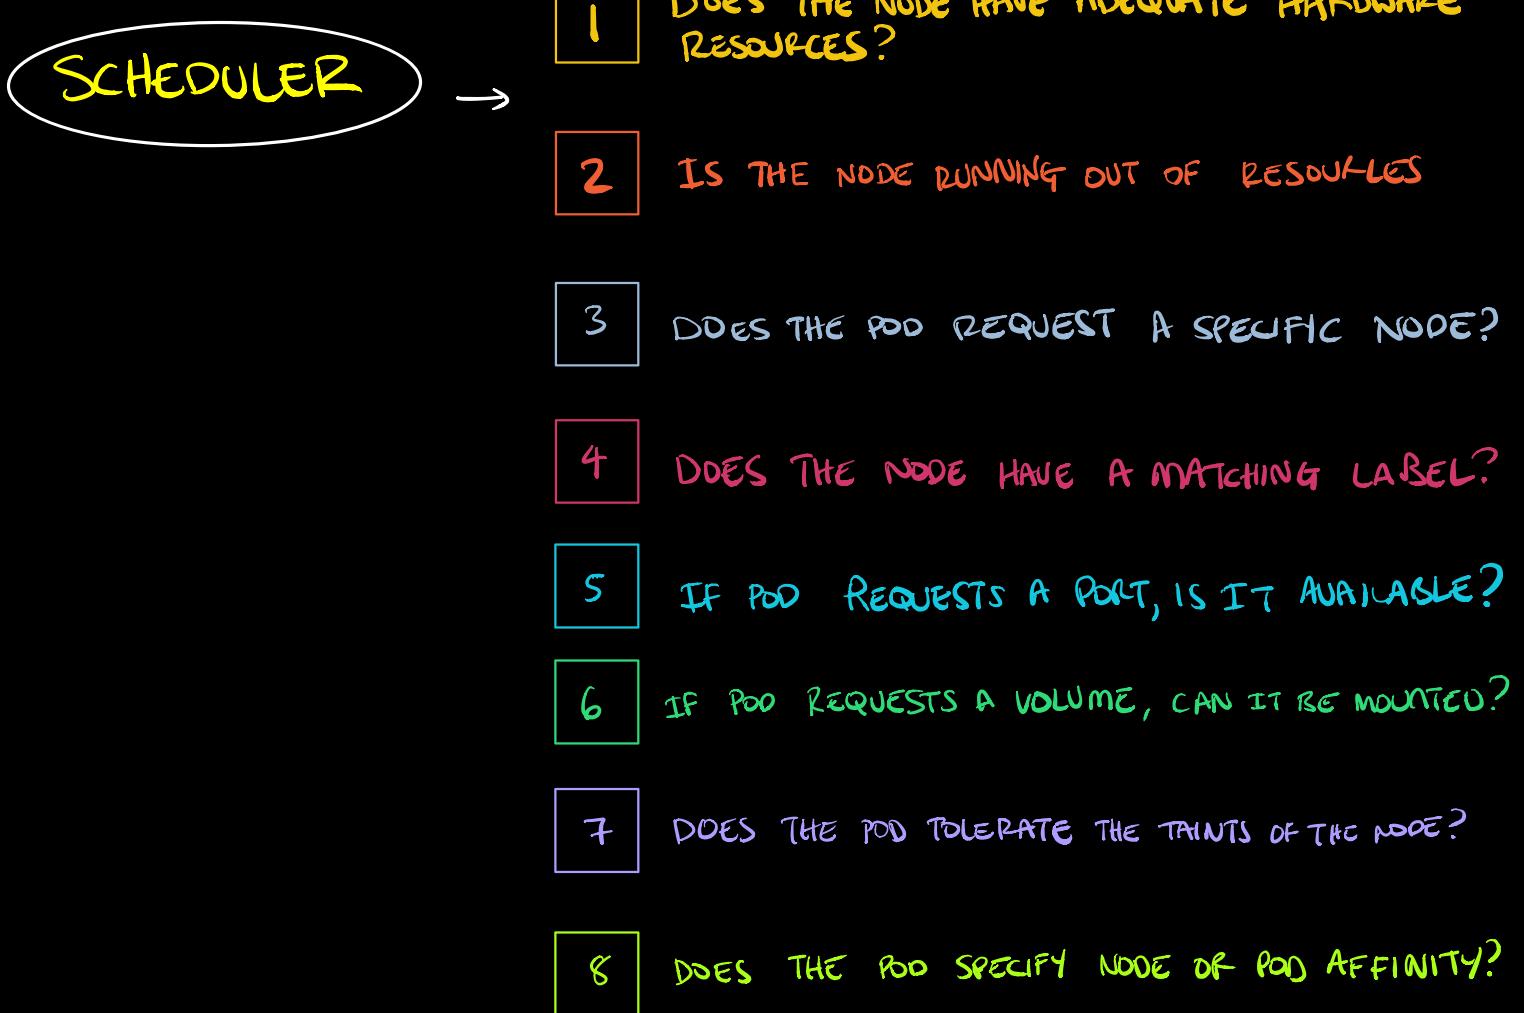


A PODS DNS SEARCH WILL INCLUDE THE PODS OWN NAMESPACE AND THE CLUSTERS DEFAULT DOMAIN

# POD SCHEDULING WITHIN THE KUBERNETES CLUSTER

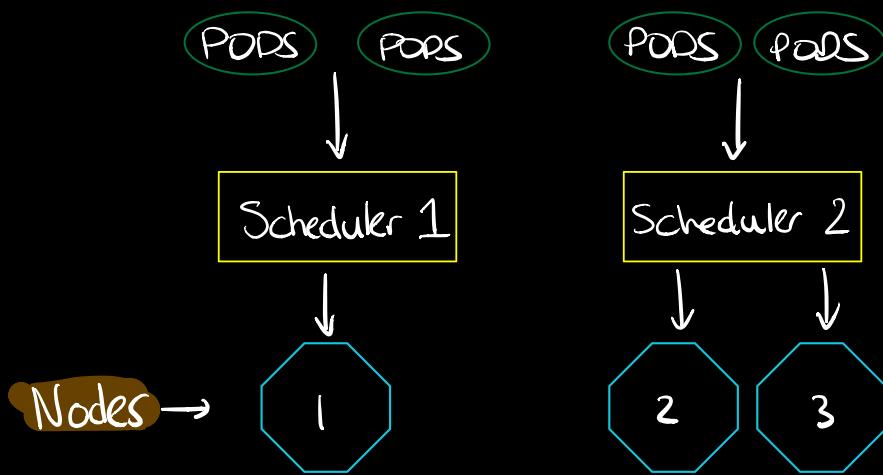
# CONFIGURING THE KUBERNETES SCHEDULER

SCHEDULER RESPONSIBLE FOR ASSIGNING POD TO NODE  
BASED ON RESOURCE REQUIREMENTS OF THE POD



# RUNNING MULTIPLE SCHEDULERS FOR MULTIPLE PODS

It is possible to have 2 schedulers working alongside each other.

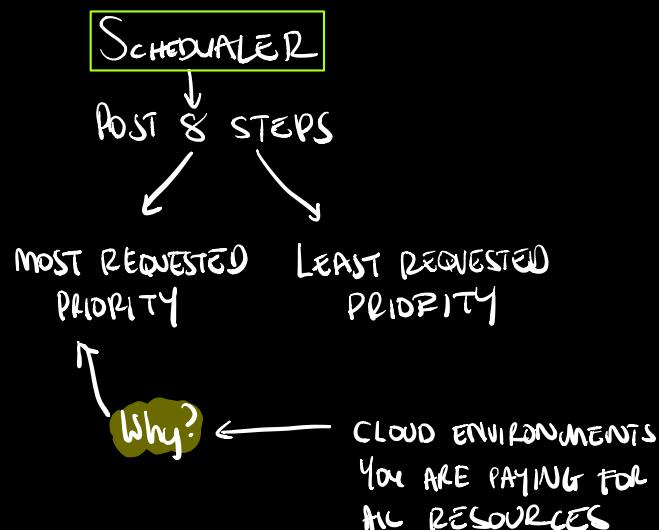


## SCHEDULING PODS WITH LIMITS AND LABEL SELECTORS

TRAINTS → REPEL WORK → EXAMPLE  
MASTER NODE  
NO SCHEDUAL

TOLERATIONS → ALLOW YOU TO  
TOLERATE A TRAINT → EXAMPLE  
KUBE-PROXY ← DAEMON SET POD  
MUST RUN ON ALL  
NODES

CPU/MEMORY  
↓  
POD MAY NOT BE  
USING ALL REQUESTED  
RESOURCE AT A GIVEN  
TIME.  
SCHEDULER LOOKS  
AT THE SUM OF  
RESOURCES REQUESTED  
BY EXISTING PODS



# DAEMONSETS

DaemonSets ensure that a single replica of a pod is running on each node at all times



POD DAEMONSET POD  
POD REPLICASET POD

If you try delete a daemonset pod, it will simply recreate it.

# DISPLAY SCHEDULER EVENTS

