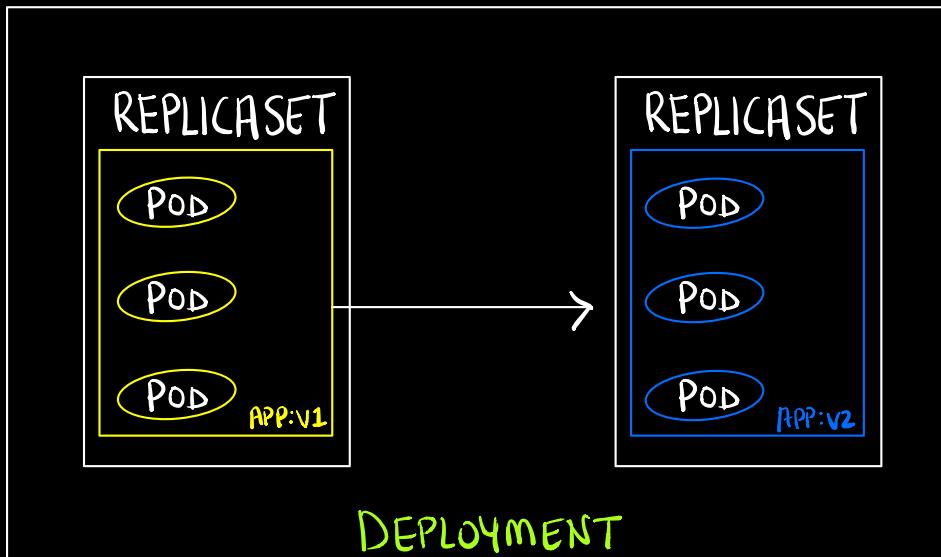


DEPLOYING
APPLICATIONS
IN THE
KUBERNETES
CLUSTER

DEPLOYING AN APPLICATION, ROLLING UPDATES, AND ROLLBACKS

DEPLOYMENTS → HIGH LEVEL RESOURCE FOR
DEPLOYING AND UPDATING APPS



KUBECTL APPLY → MODIFY OBJECTS TO EXISTING YAML
IF DEPLOYMENT NOT CREATED → ALSO CREATE

KUBECTL REPLACE → REPLACES OLD WITH NEW AND OBJECT MUST
EXIST.

ROLLING UPDATE → PREFERRED WAY → SERVICE NOT INTERRUPTED
→ FASTEST WAY

KUBECTL ROLLOUT → ROLL BACK PREVIOUS VERSION

CONFIGURING AN APP FOR HA AND SCALE

AVOIDING
BAD
VERSIONS

→ BLOCK BAD VERSION
RELEASE

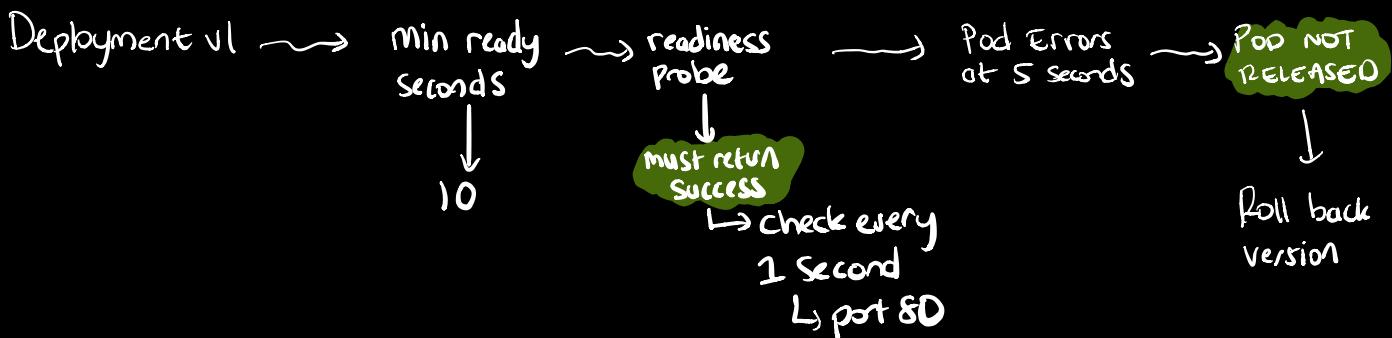
MIN READY
SECONDS

How long a
newly created
Pod should be
ready before
considered
available



READINESS
PROBE

Determines
if a specific
Pod should
receive
client request
or not.



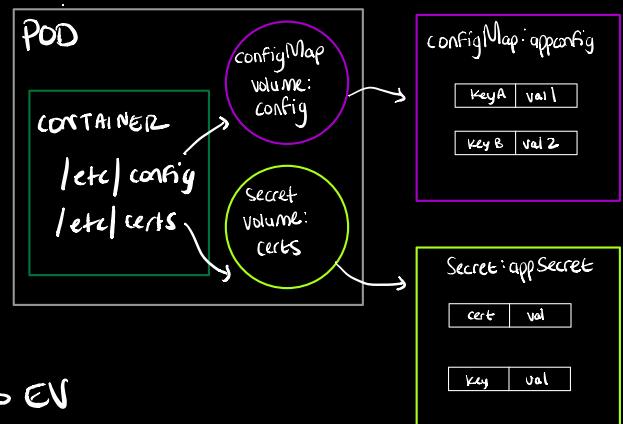
PASSING CONFIGURATION OPTIONS TO APP

ENVIRONMENT VARIABLES

STORE IN CONFIG MAP
CREATE SECRET & PASS TO EV

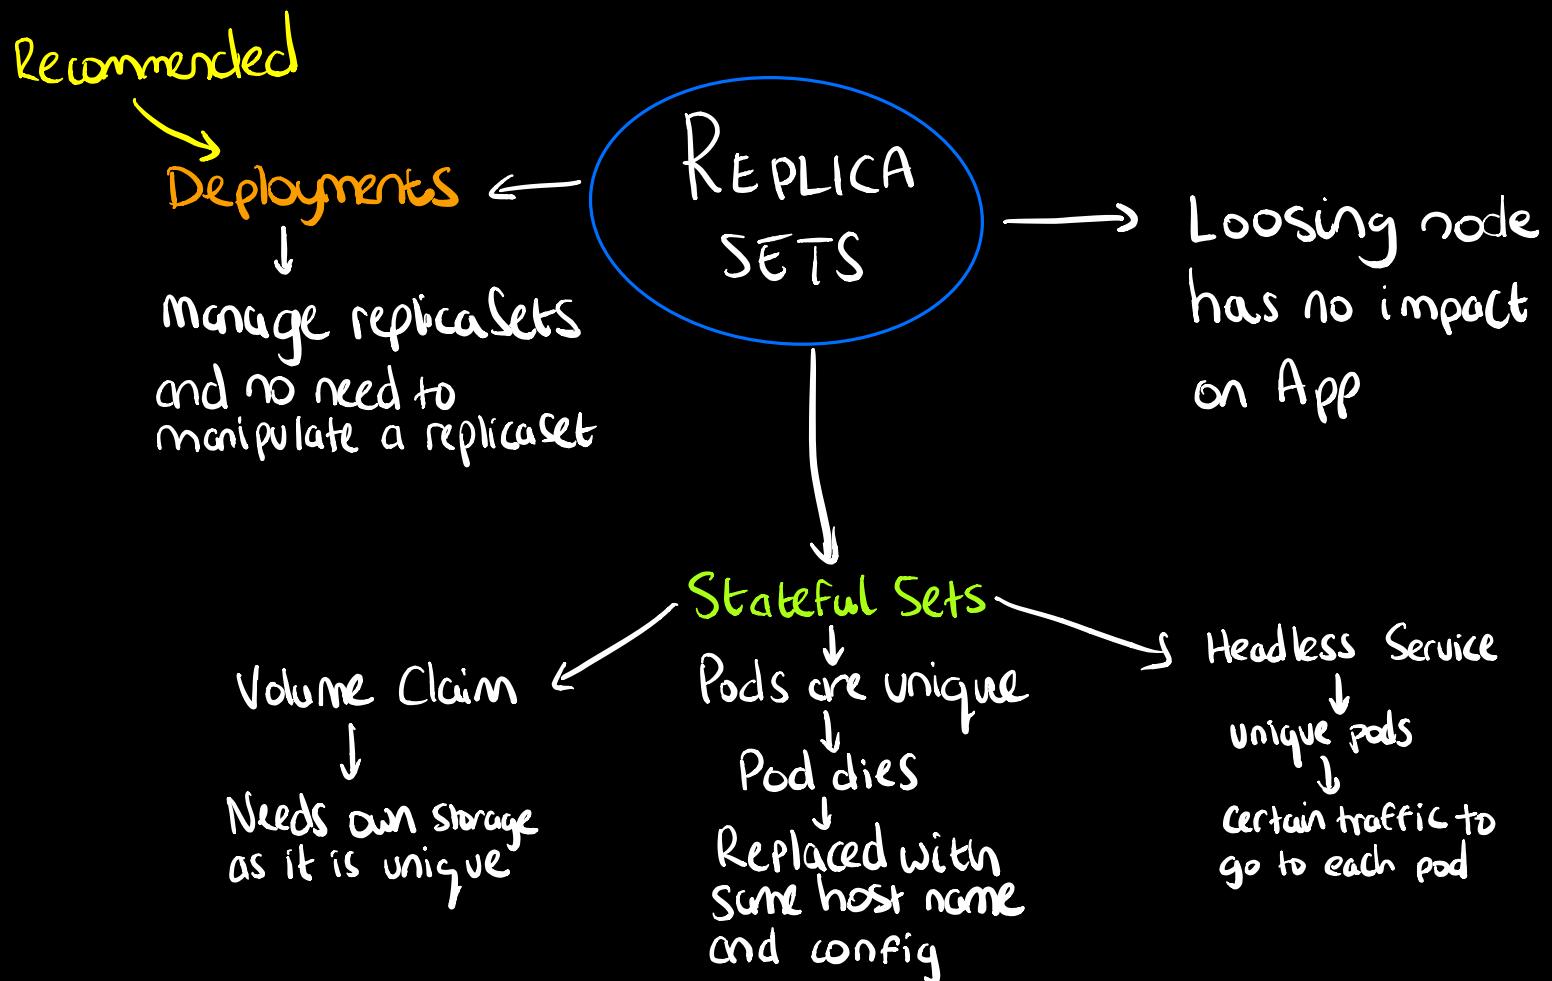
JUST UPDATE → NO NEED TO REBUILD
IMAGE

MULTIPLE
CONTAINERS
CAN USE
SAME



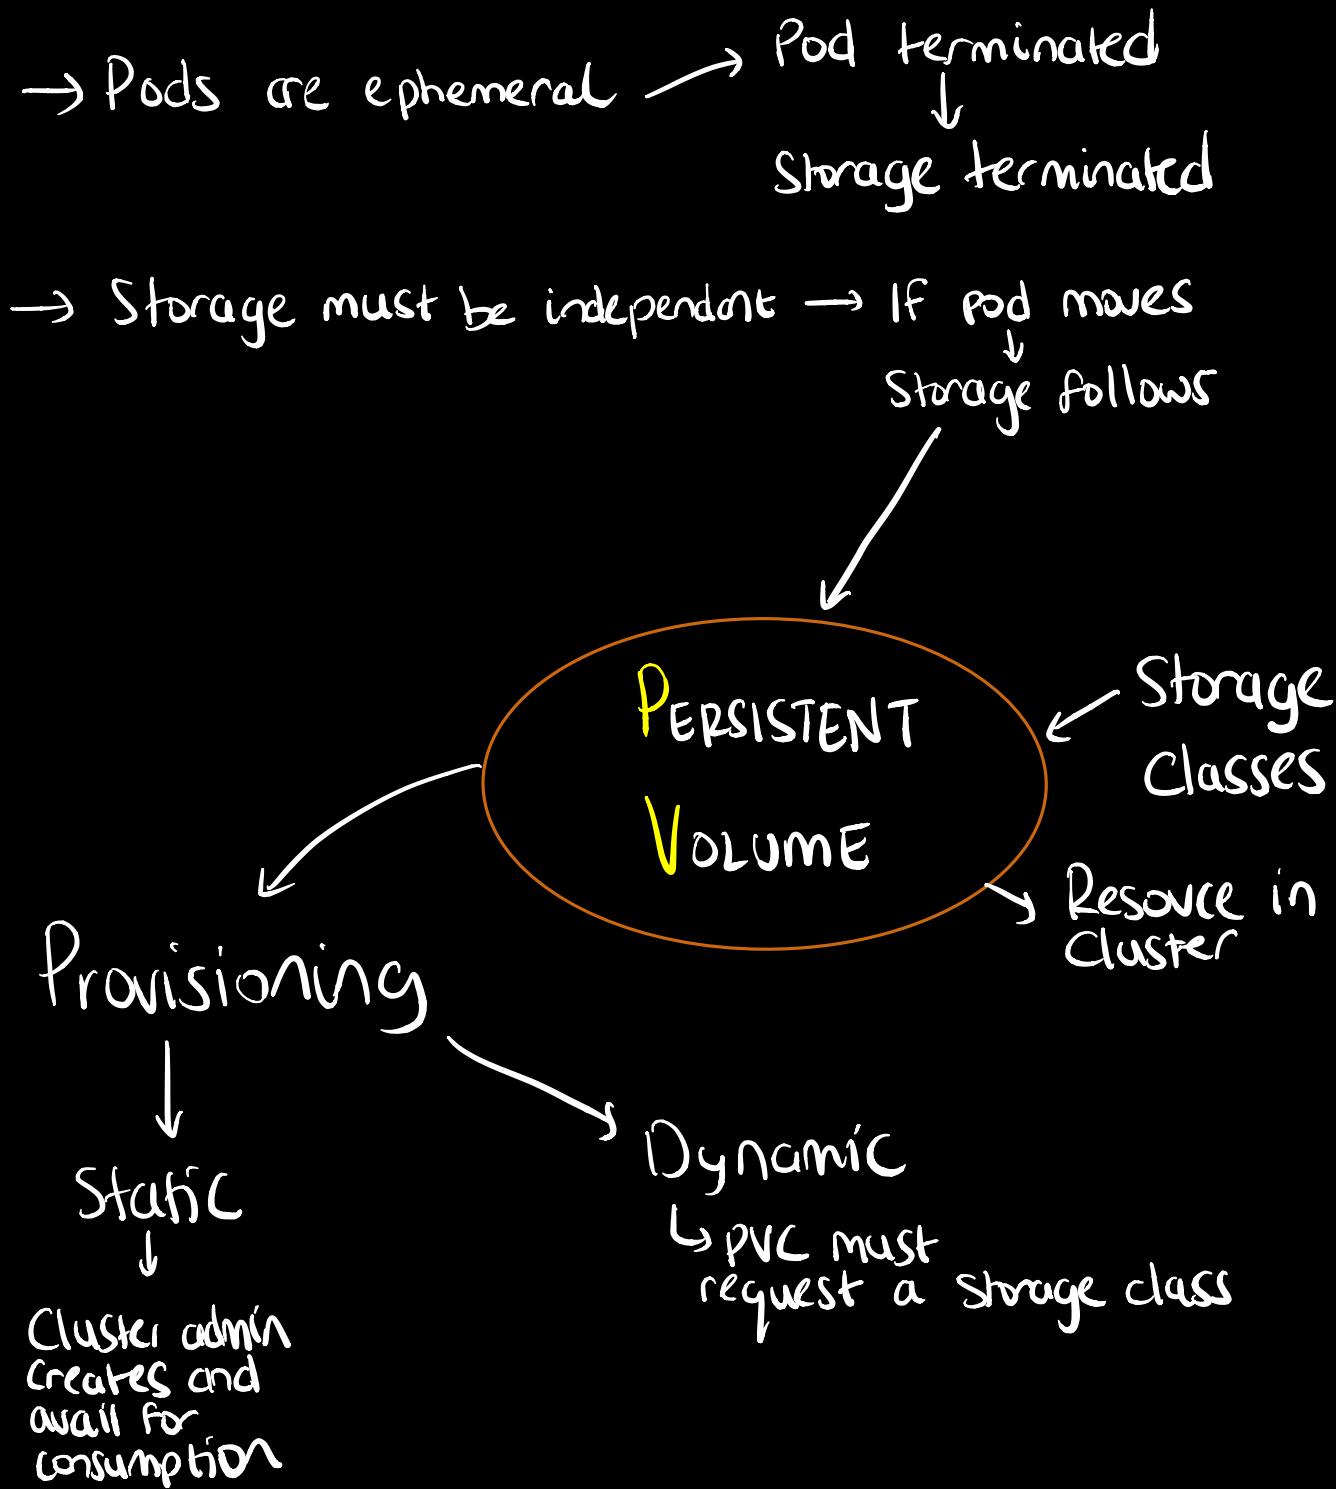
CREATING A SELF-HEALING APP

ReplicaSets ensure that a identically configured pods are running at the desired replica count



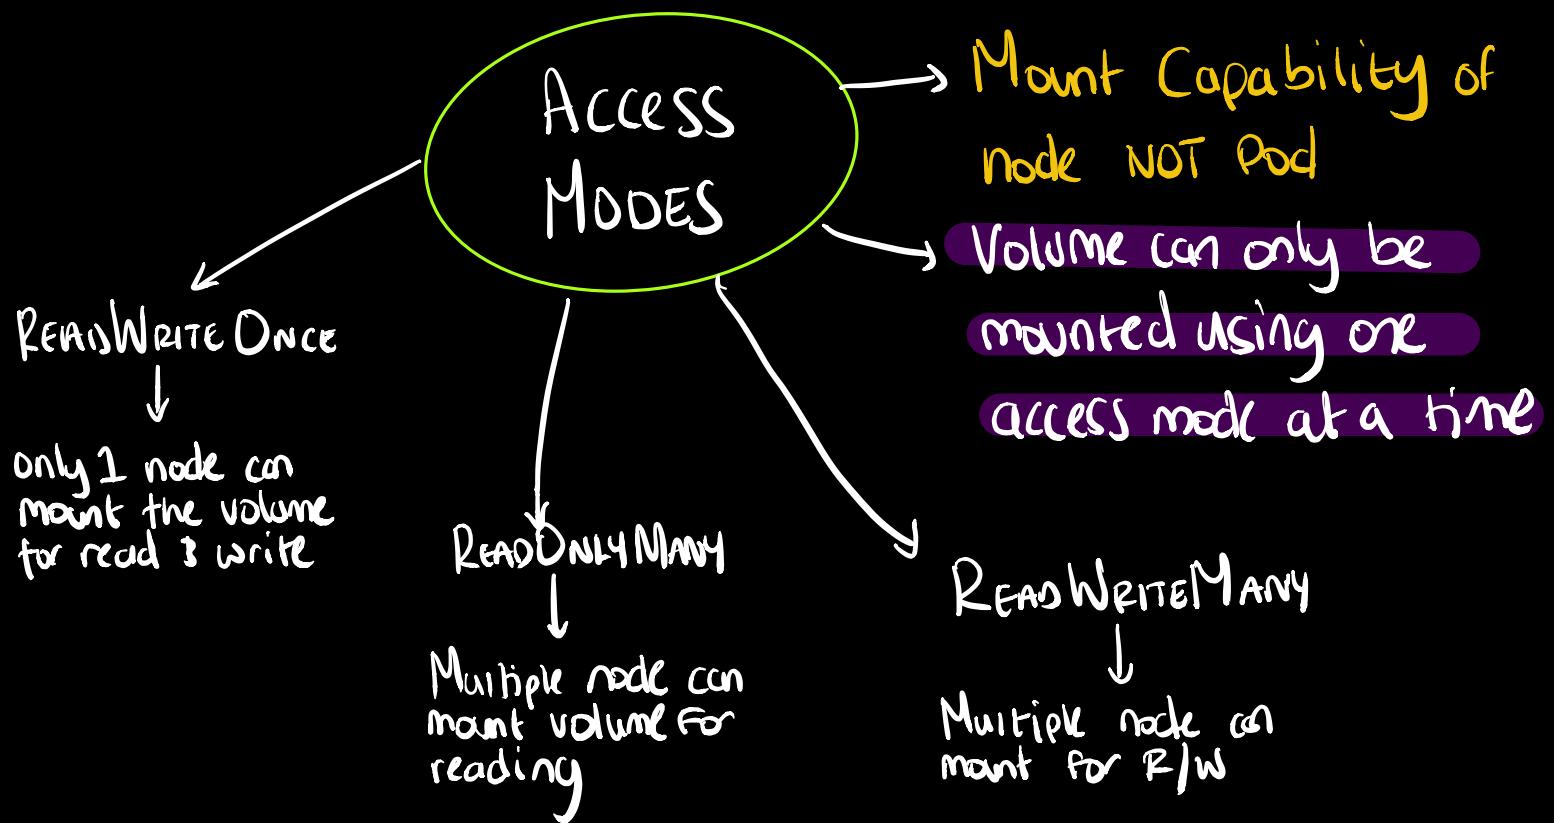
MANAGING DATA IN KUBERNETES CLUSTER

PERSISTENT VOLUMES



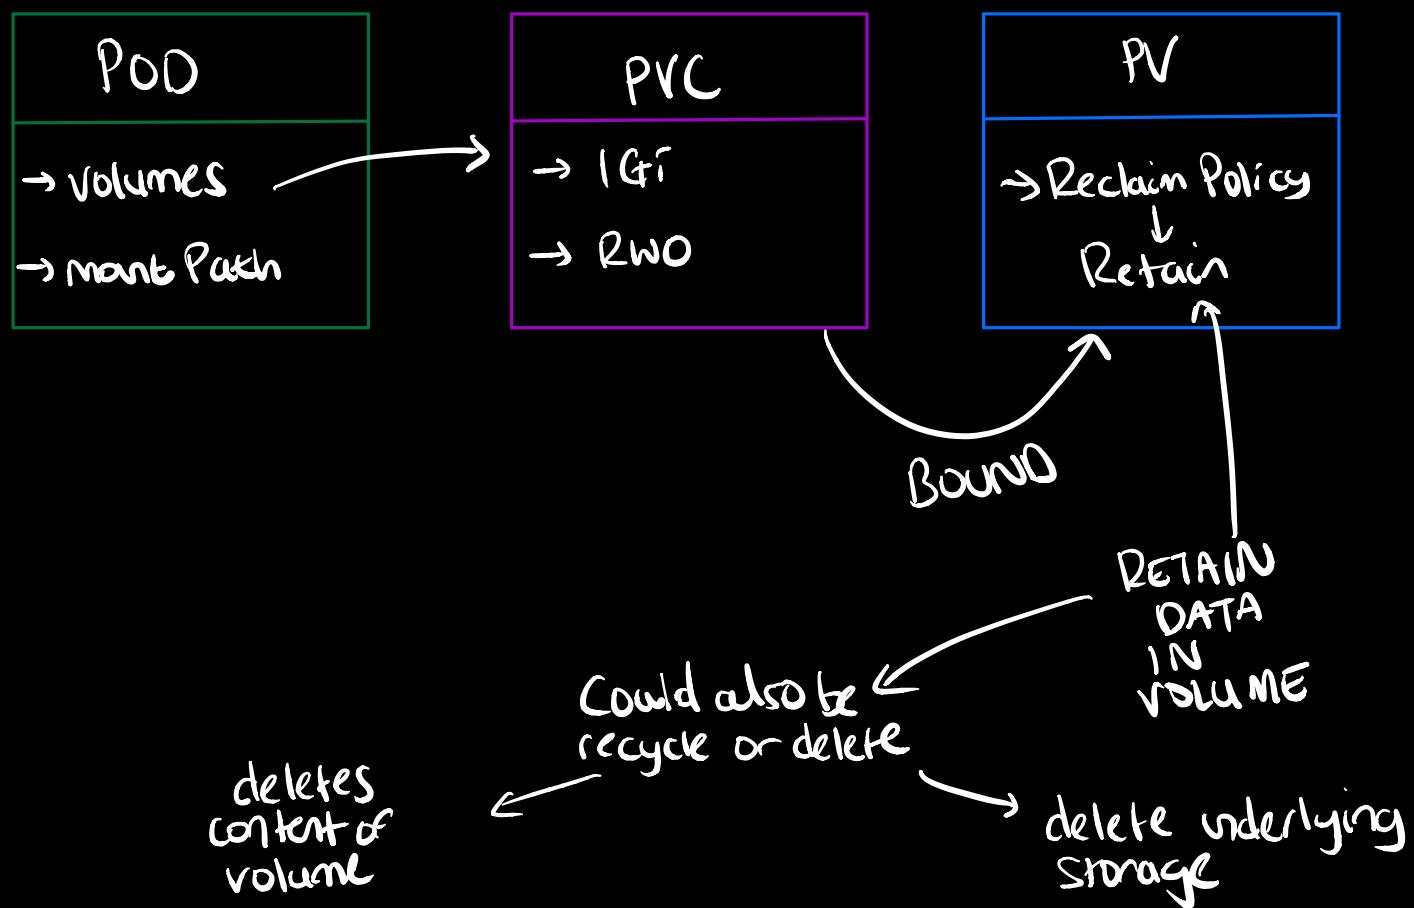
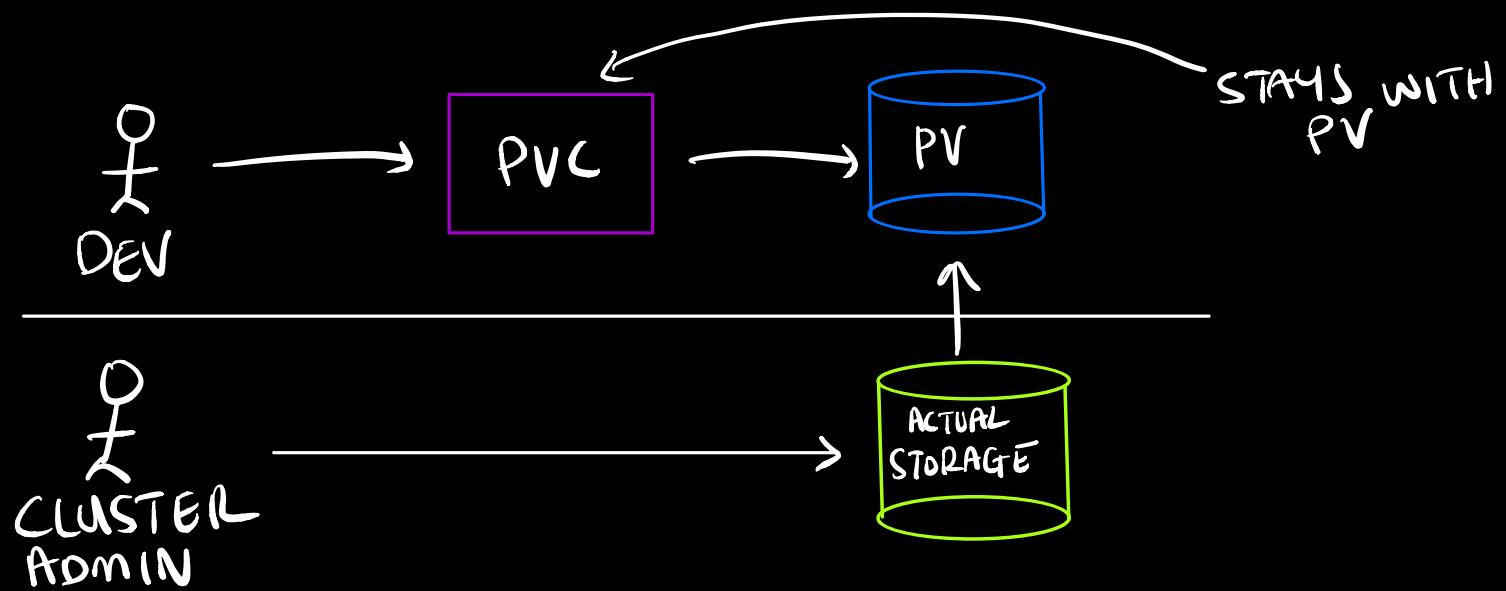
VOLUME ACCESS MODES

By specifying an access mode with your PV,
you allow the volume to be mounted to one
or many nodes, as well as read by one or many



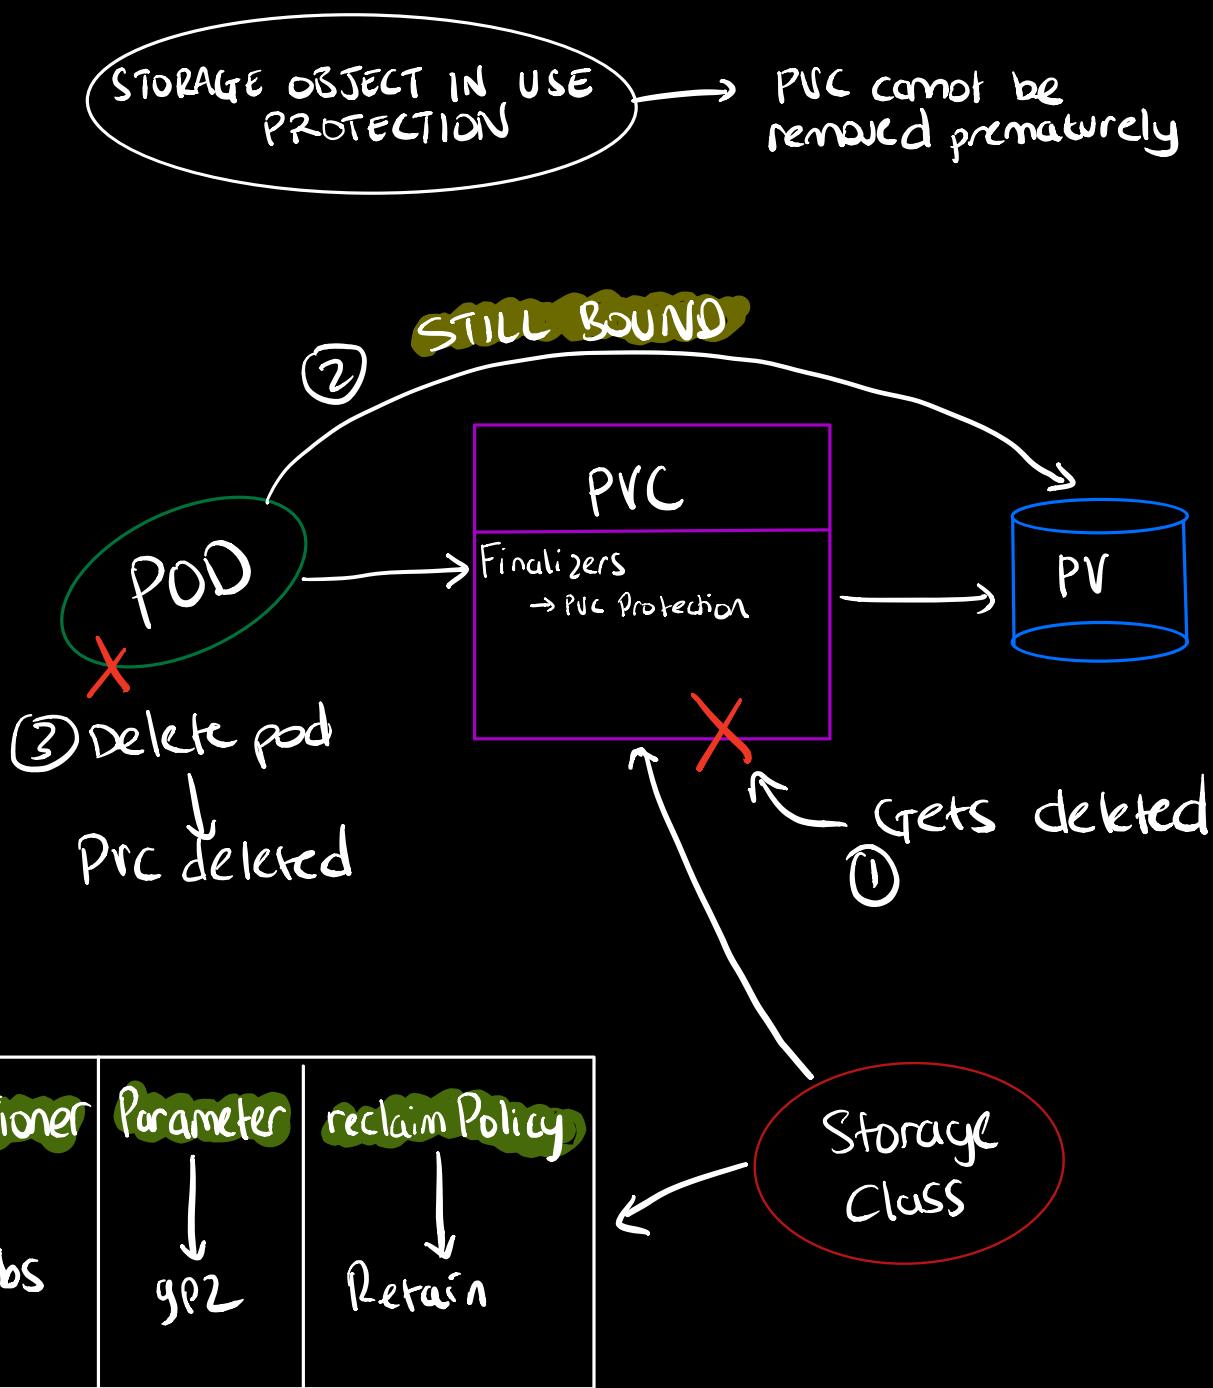
PERSISTENT VOLUME CLAIMS (PVC)

PVC allows the application developer to request storage for the application, without having to know underlying infra.



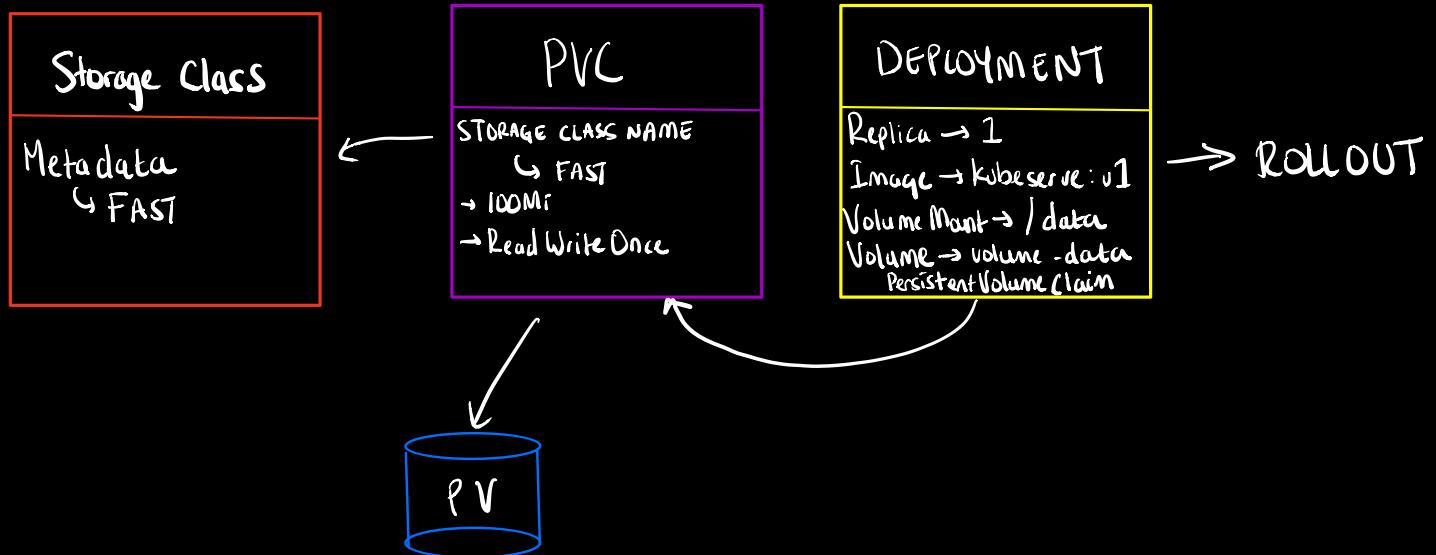
STORAGE OBJECTS

VOLUMES that are already in use by a pod are protected against data loss. This means even if you delete a PVC, you can still access volume from pod.



APPLICATIONS WITH PERSISTENT STORAGE

Example →

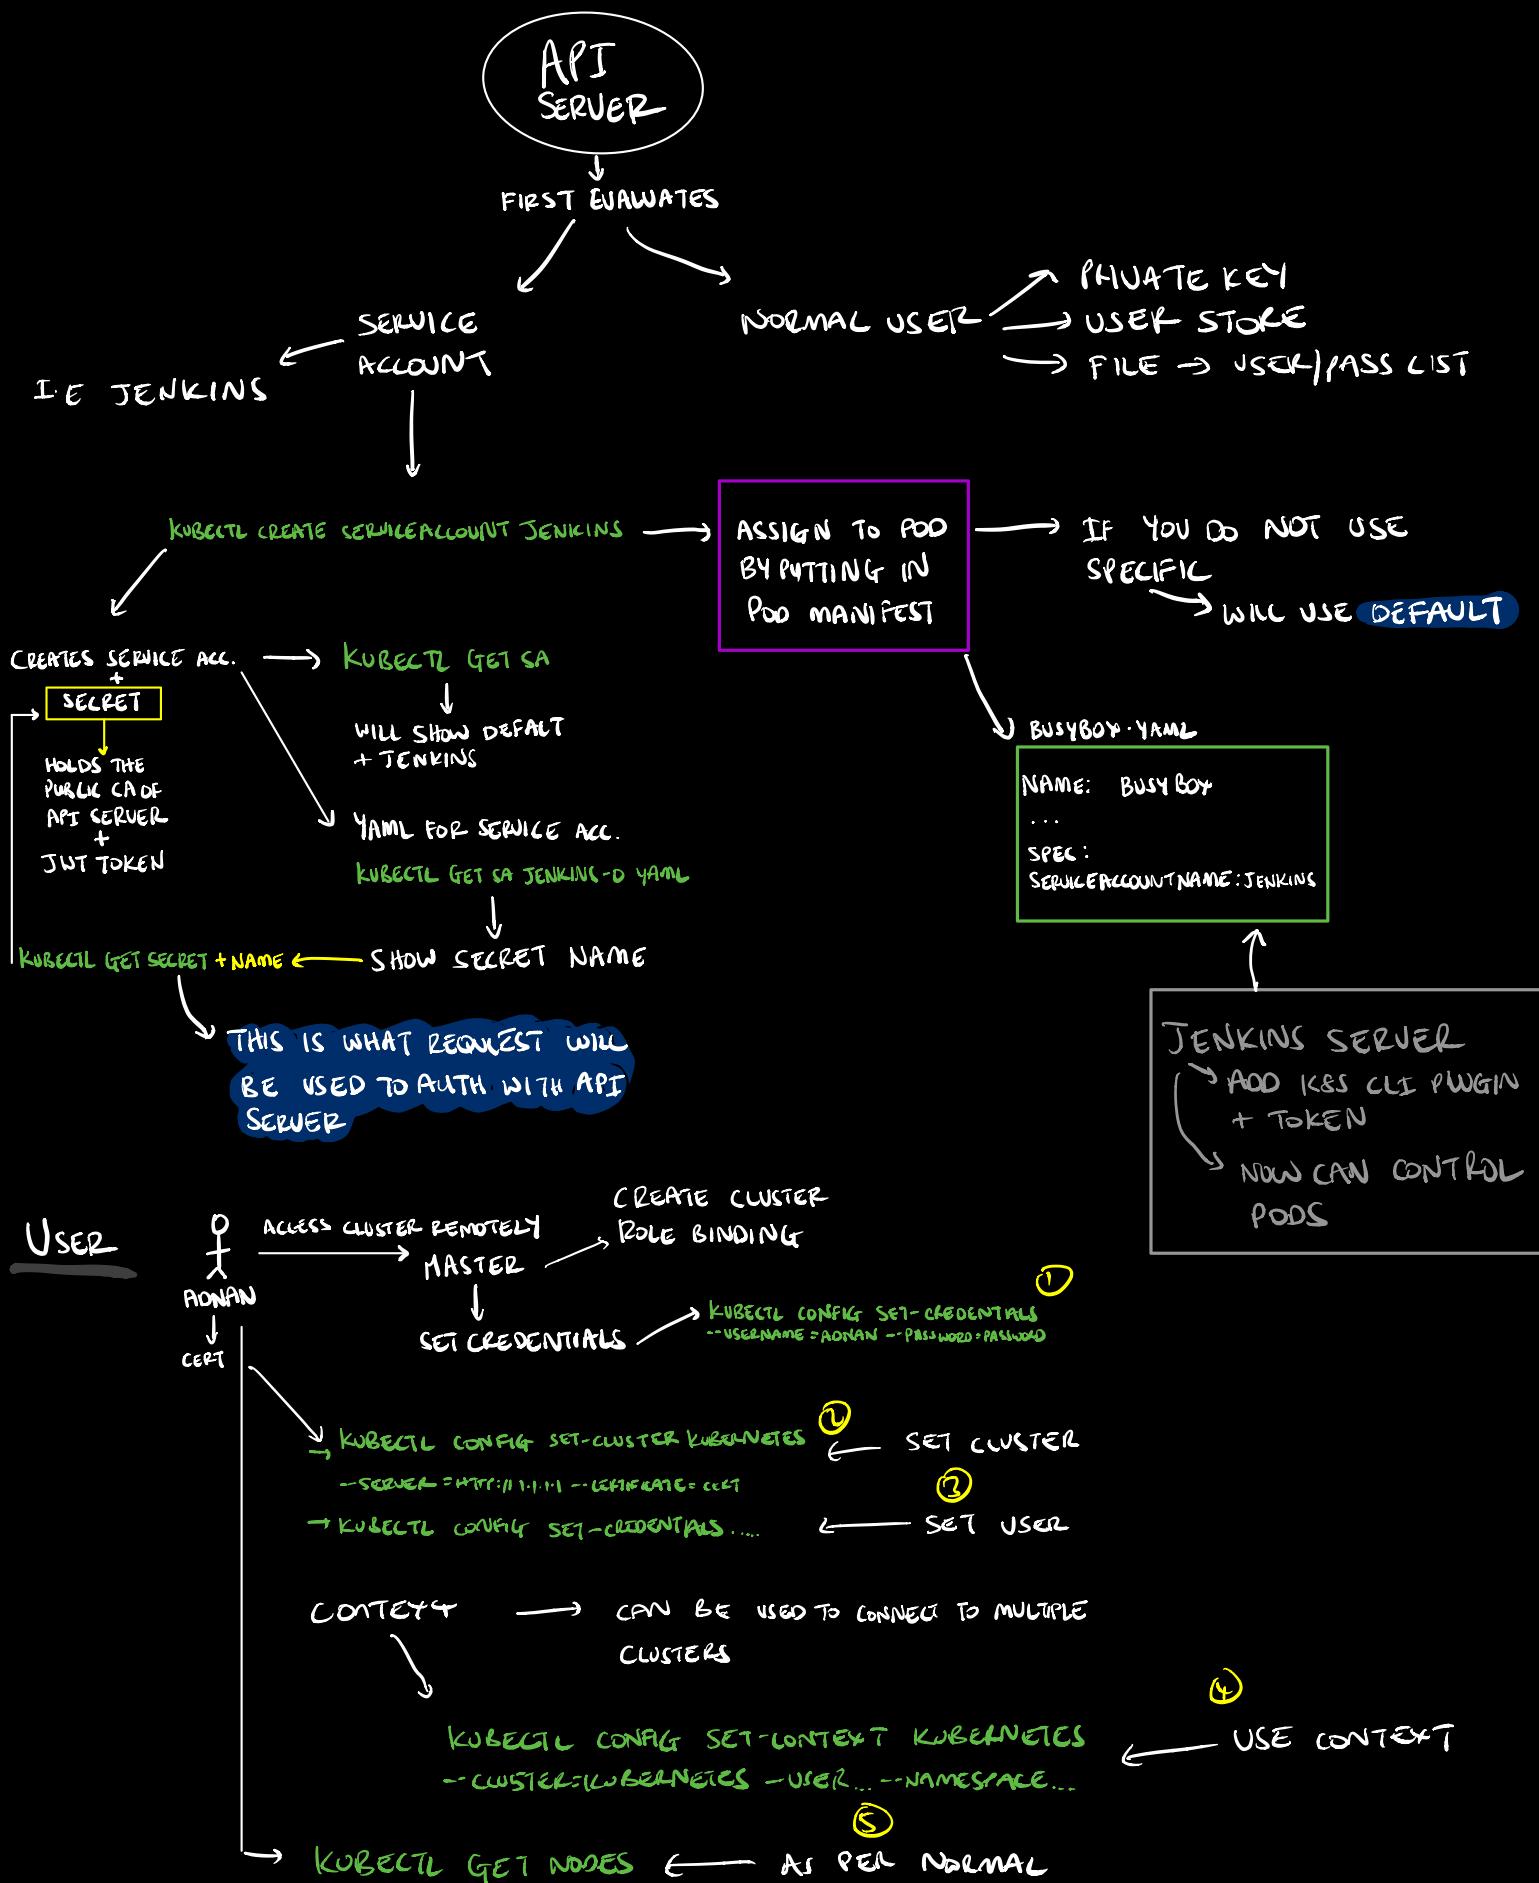


- ① create storage class object
- ② create PVC object
- ③ create deployment
- ④ rollout deployment
- ⑤ check pods
- ⑥ create file on mount
- ⑦ List contents.

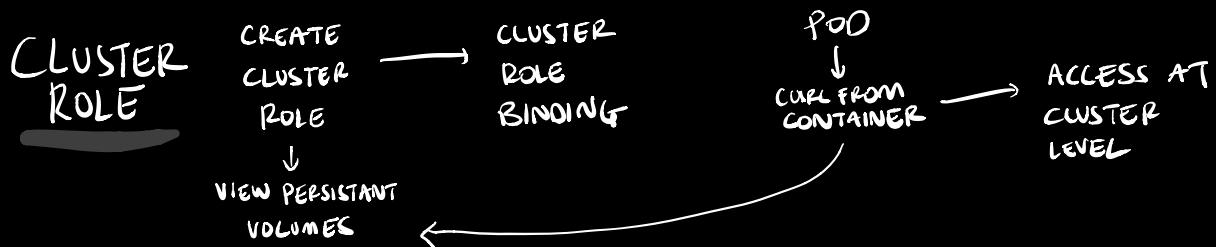
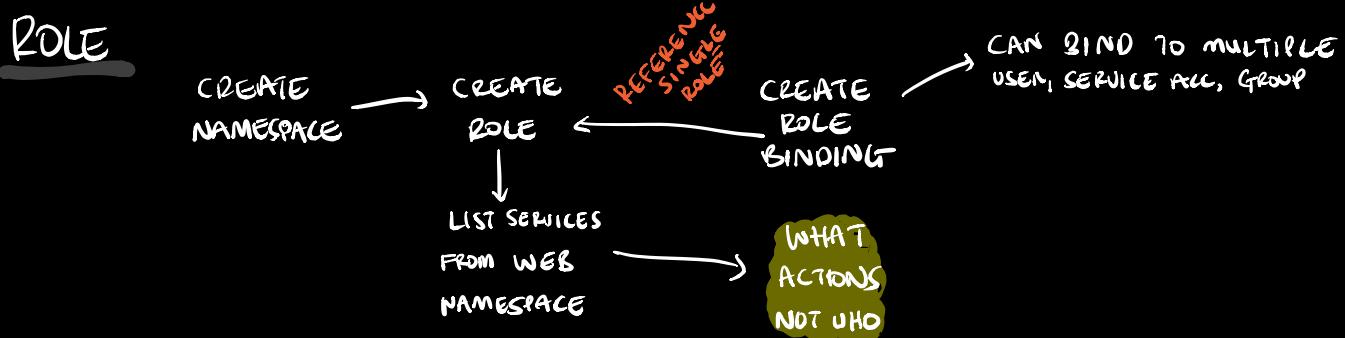
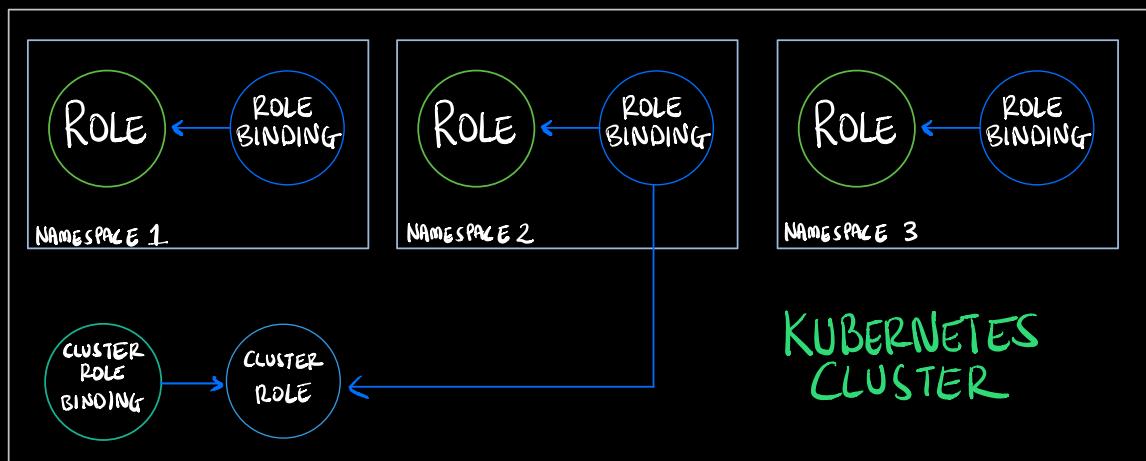
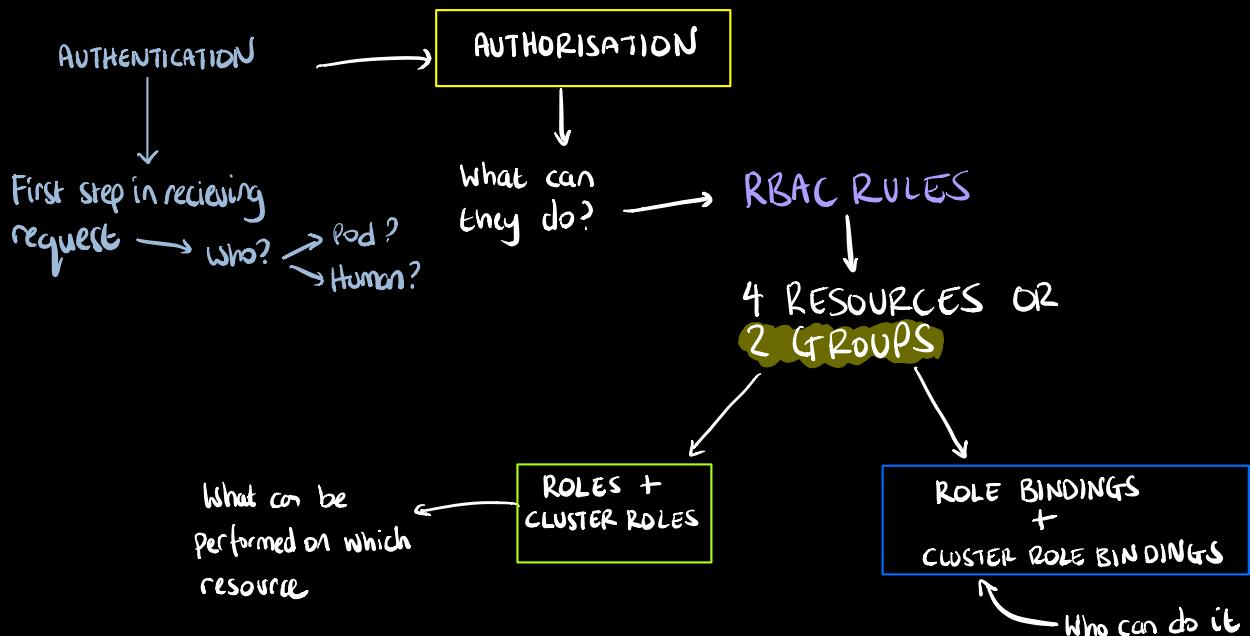


SECURING THE KUBERNETES CLUSTER

SERVICE ACCOUNTS AND USERS

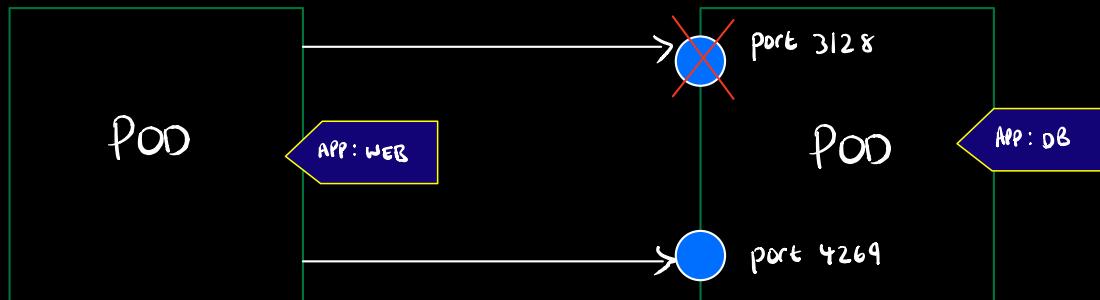


CLUSTER AUTHENTICATION AND AUTHORISATION

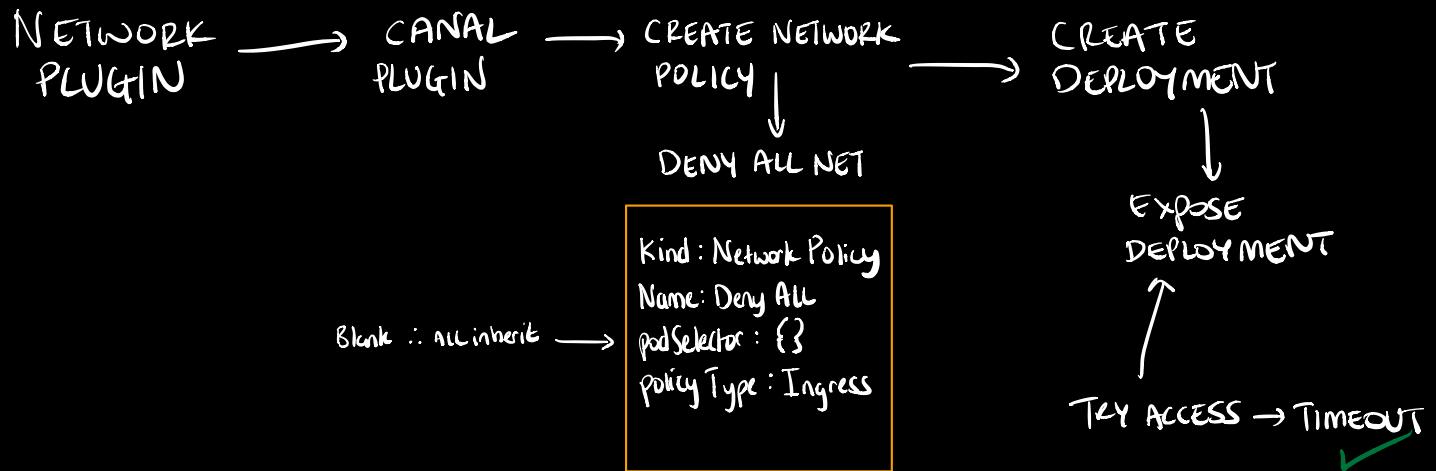


CONFIGURING NETWORK POLICIES

Network policies use selectors to apply rules to pods for communication throughout the cluster



How?



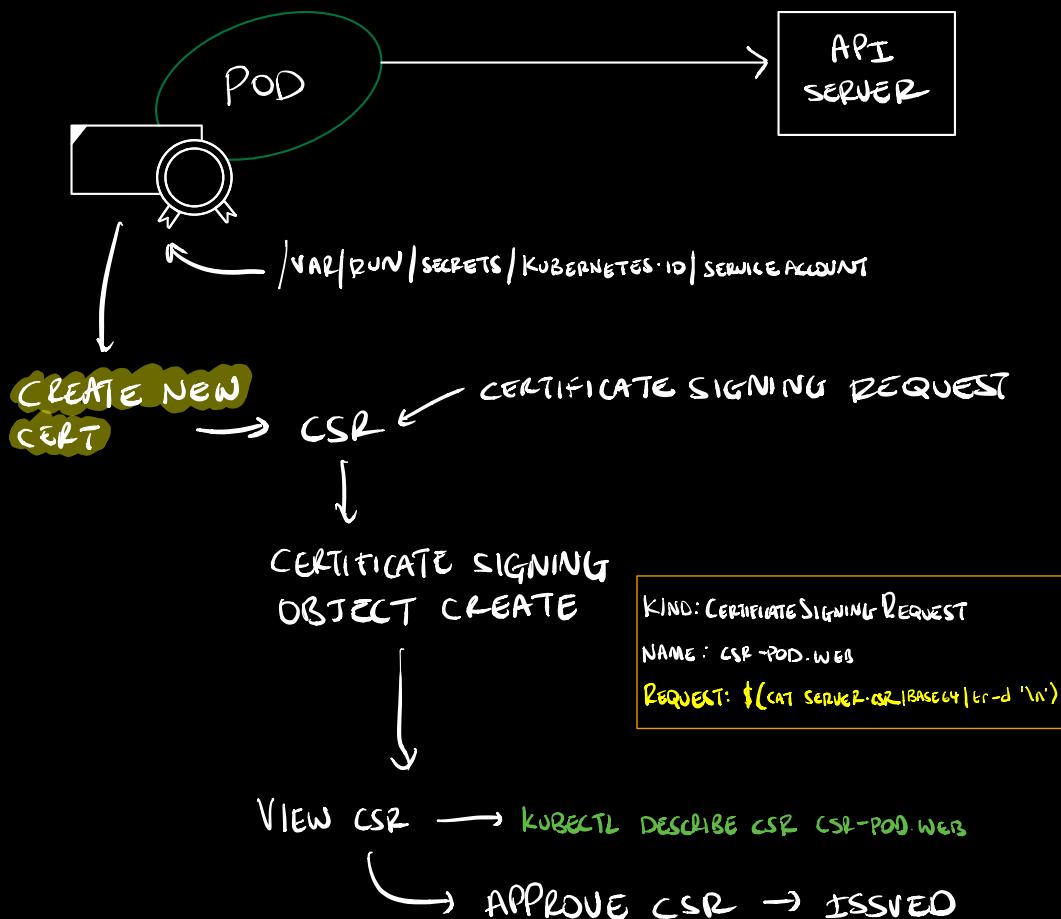
→ CREATE NETWORK POLICY → LABEL PODS

podSelector:
matchLabels:
 app: db
ingress:
 ...
 matchLabels
 app: web
port:
 ~port: 4269

Allow comms
between pods
and specified
port

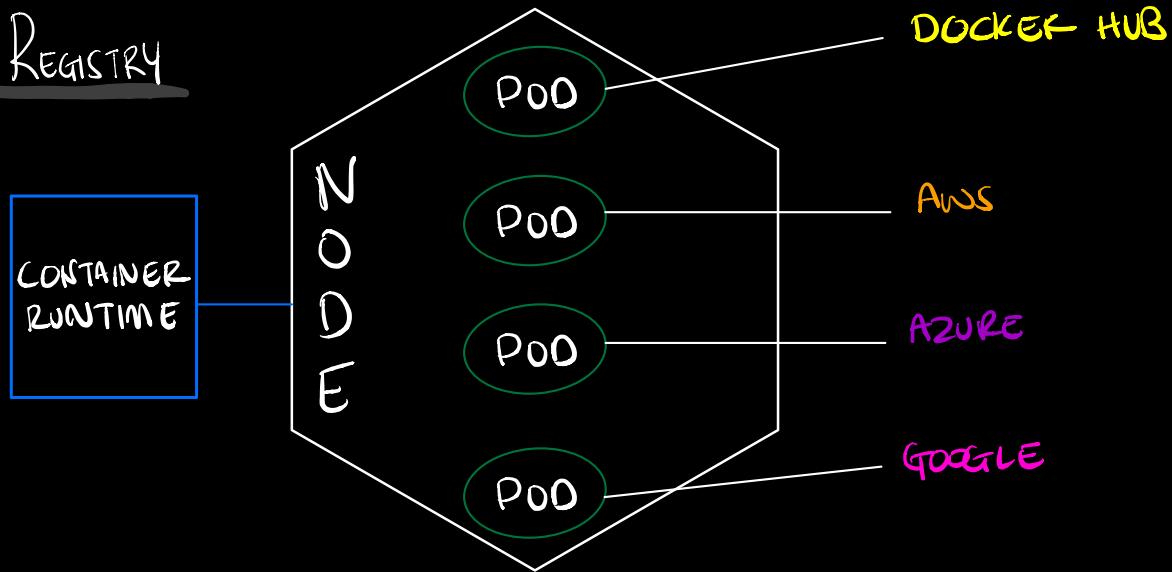
CREATING TLS CERTIFICATES

The CA is used to generate a TLS certificate and authenticate with the API Server.



SECURE IMAGES

PRIVATE REGISTRY



CONTROLLING
IMAGES THAT
GO INTO
PRODUCTION

VULNERABILITIES → CLAIR (SCANNING)

SOMETHING ON IT
CAUSE NODE CRASH

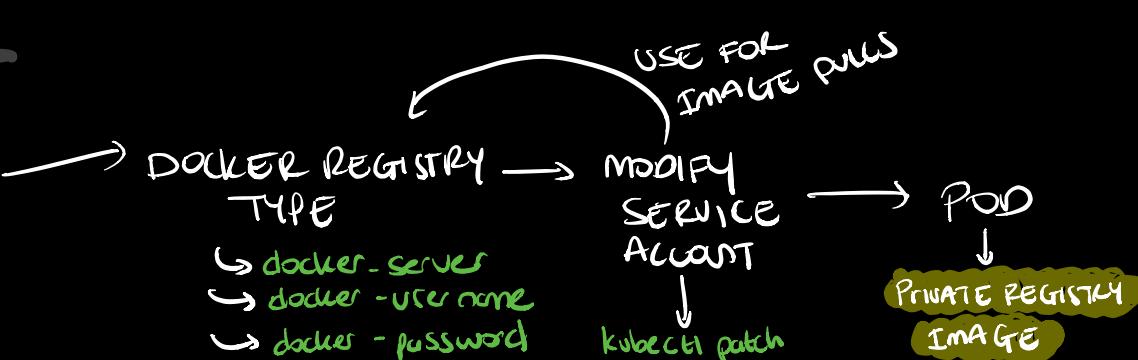
LOGIN TO PRIVATE REGISTRY

↳ TAG DOCKER IMAGE

↳ PUSH TO PRIVATE REGISTRY

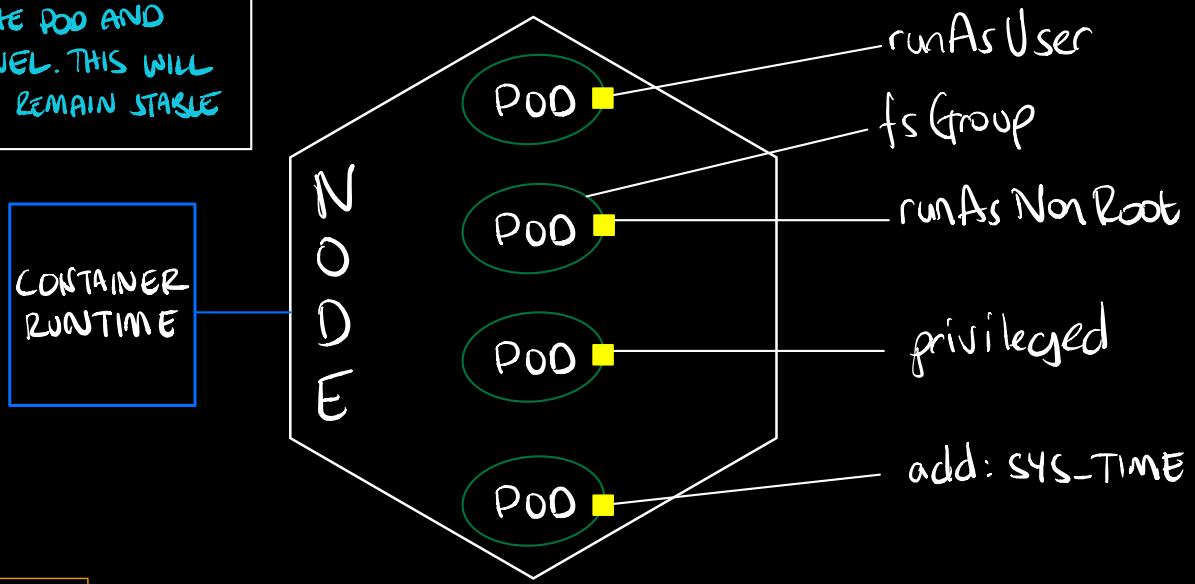
How?

KUBERNETES
CREATE
SECRET



DEFINING SECURITY CONTEXT

LIMIT ACCESS TO CERTAIN OBJECTS AT THE POD AND CONTAINER LEVEL. THIS WILL ALLOW IMAGES TO REMAIN STABLE



Kind: Pod
image: alpine
Security Context:
runAsUser: 405

Run Pod as 405

Can also put 'runAsNonRoot'

Ability to run as privileged → 'privileged: true'

CONTAINER LEVEL

ABILITY TO LOCK DOWN KERNEL LEVEL FEATURES ON CONTAINER

SETTING CAPABILITIES ON POD LEVEL

ADD

Security Context:
add:
- SYS_TIME
- NET_ADMIN

REMOVE

Security Context:
drop:
CHOWN

SECURING PERSISTENT KEY/VALUE STORE

SECRETS ALLOW YOU TO EXPOSE ENTRIES AS FILES
IN A VOLUME. KEEPING THIS DATA
SECURE IS CRITICAL TO CLUSTER SECURITY

DATA MUST
LIVE BEYOND
LIFE OF POD → SECRETS → KEY/VALUE PAIR

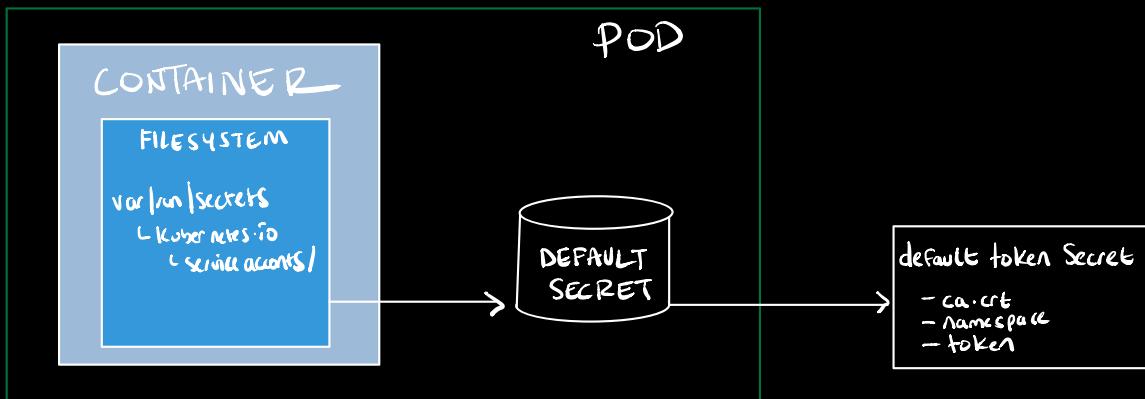
↓
PASS AS ENV VAR

OR

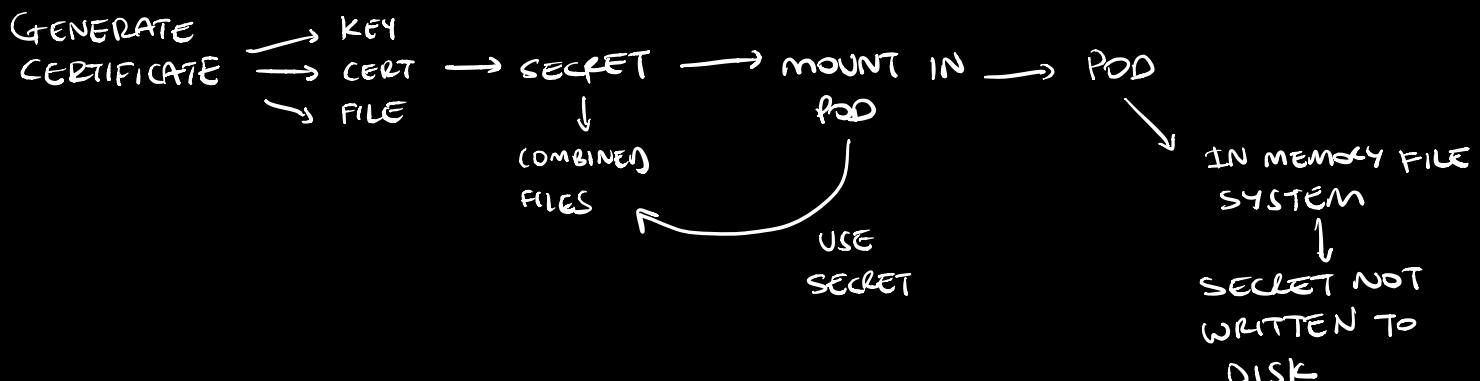
EXPOSE AS FILES
IN VOLUME

NOT BEST
PRACTICE

↓
MAY BE
OUTPUT TO
LOG FILES



HTTPS TO WEBSITE

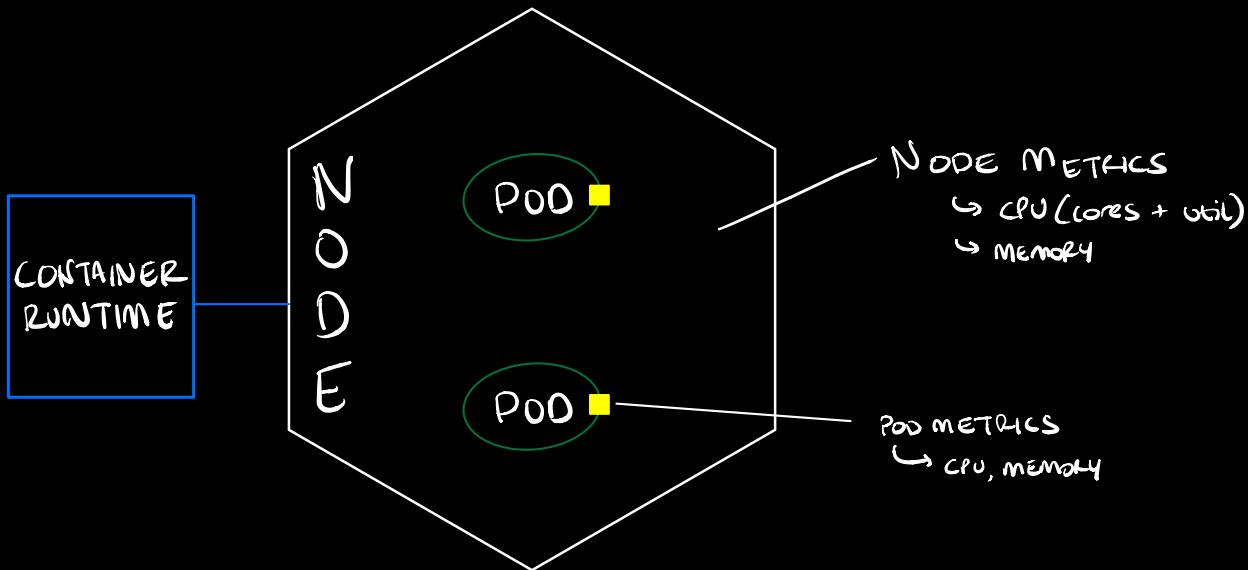




MONITORING CLUSTER COMPONENTS

MONITORING THE CLUSTER COMPONENTS

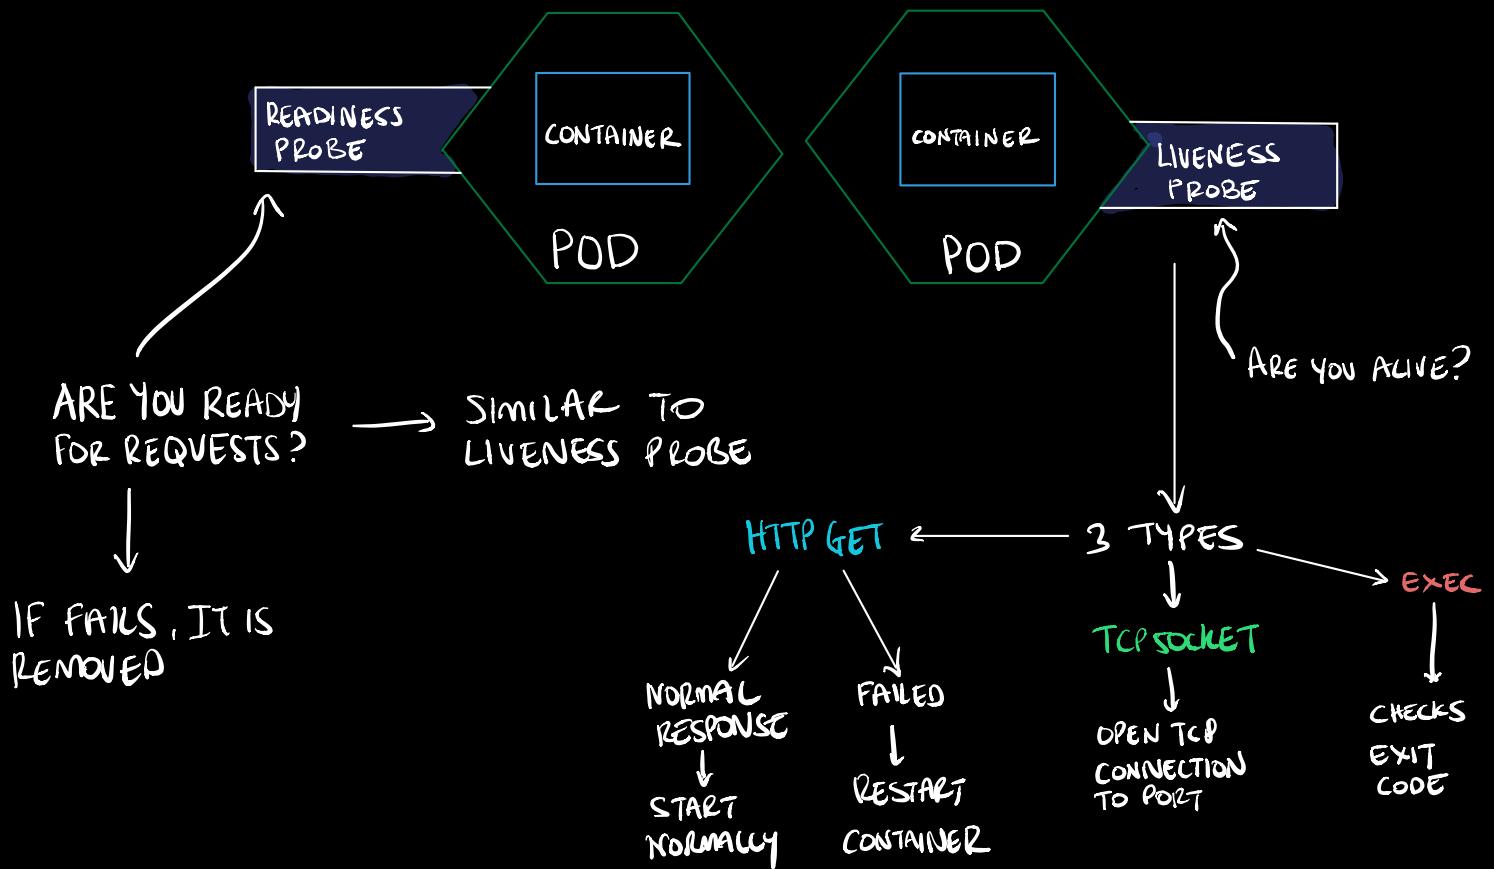
THE METRIC SERVER ALLOWS YOU TO COLLECT CPU AND MEMORY DATA FROM THE NODES AND PODS IN YOUR CLUSTER



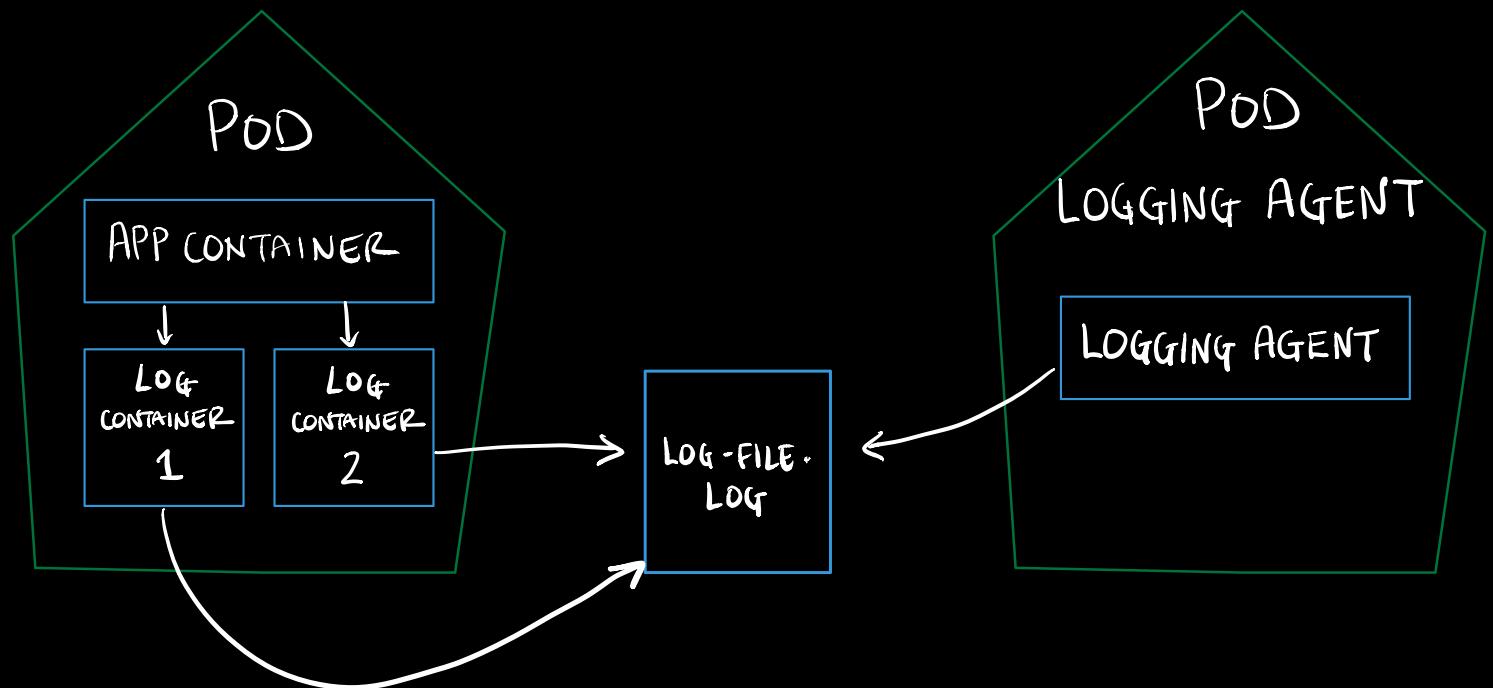
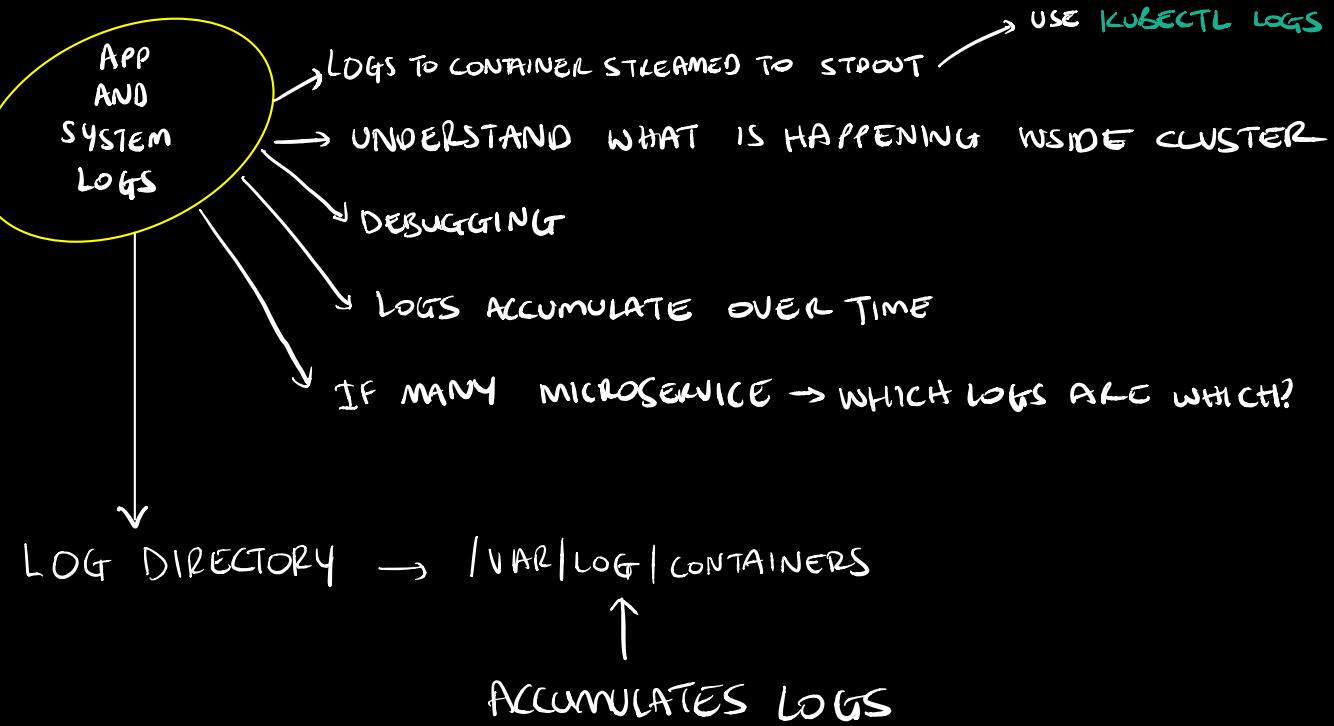
INSTALL METRIC SERVER → KUBECTL TOP NODE → CPU/MEMORY FOR ALL THE NODES
KUBECTL TOP POD → CPU/MEMORY FOR ALL THE PODS
KUBECTL TOP POD --ALL-NAMESPACES → ALL NAMESPACE
KUBECTL TOP POD -N KUBE-SYSTEM → KUBE-SYSTEM NAMESPACE
KUBECTL TOP GROUP-CONTEXT --CONTAINERS → PODS CONTAINERS

MONITORING THE APPS RUNNING WITHIN A CLUSTER

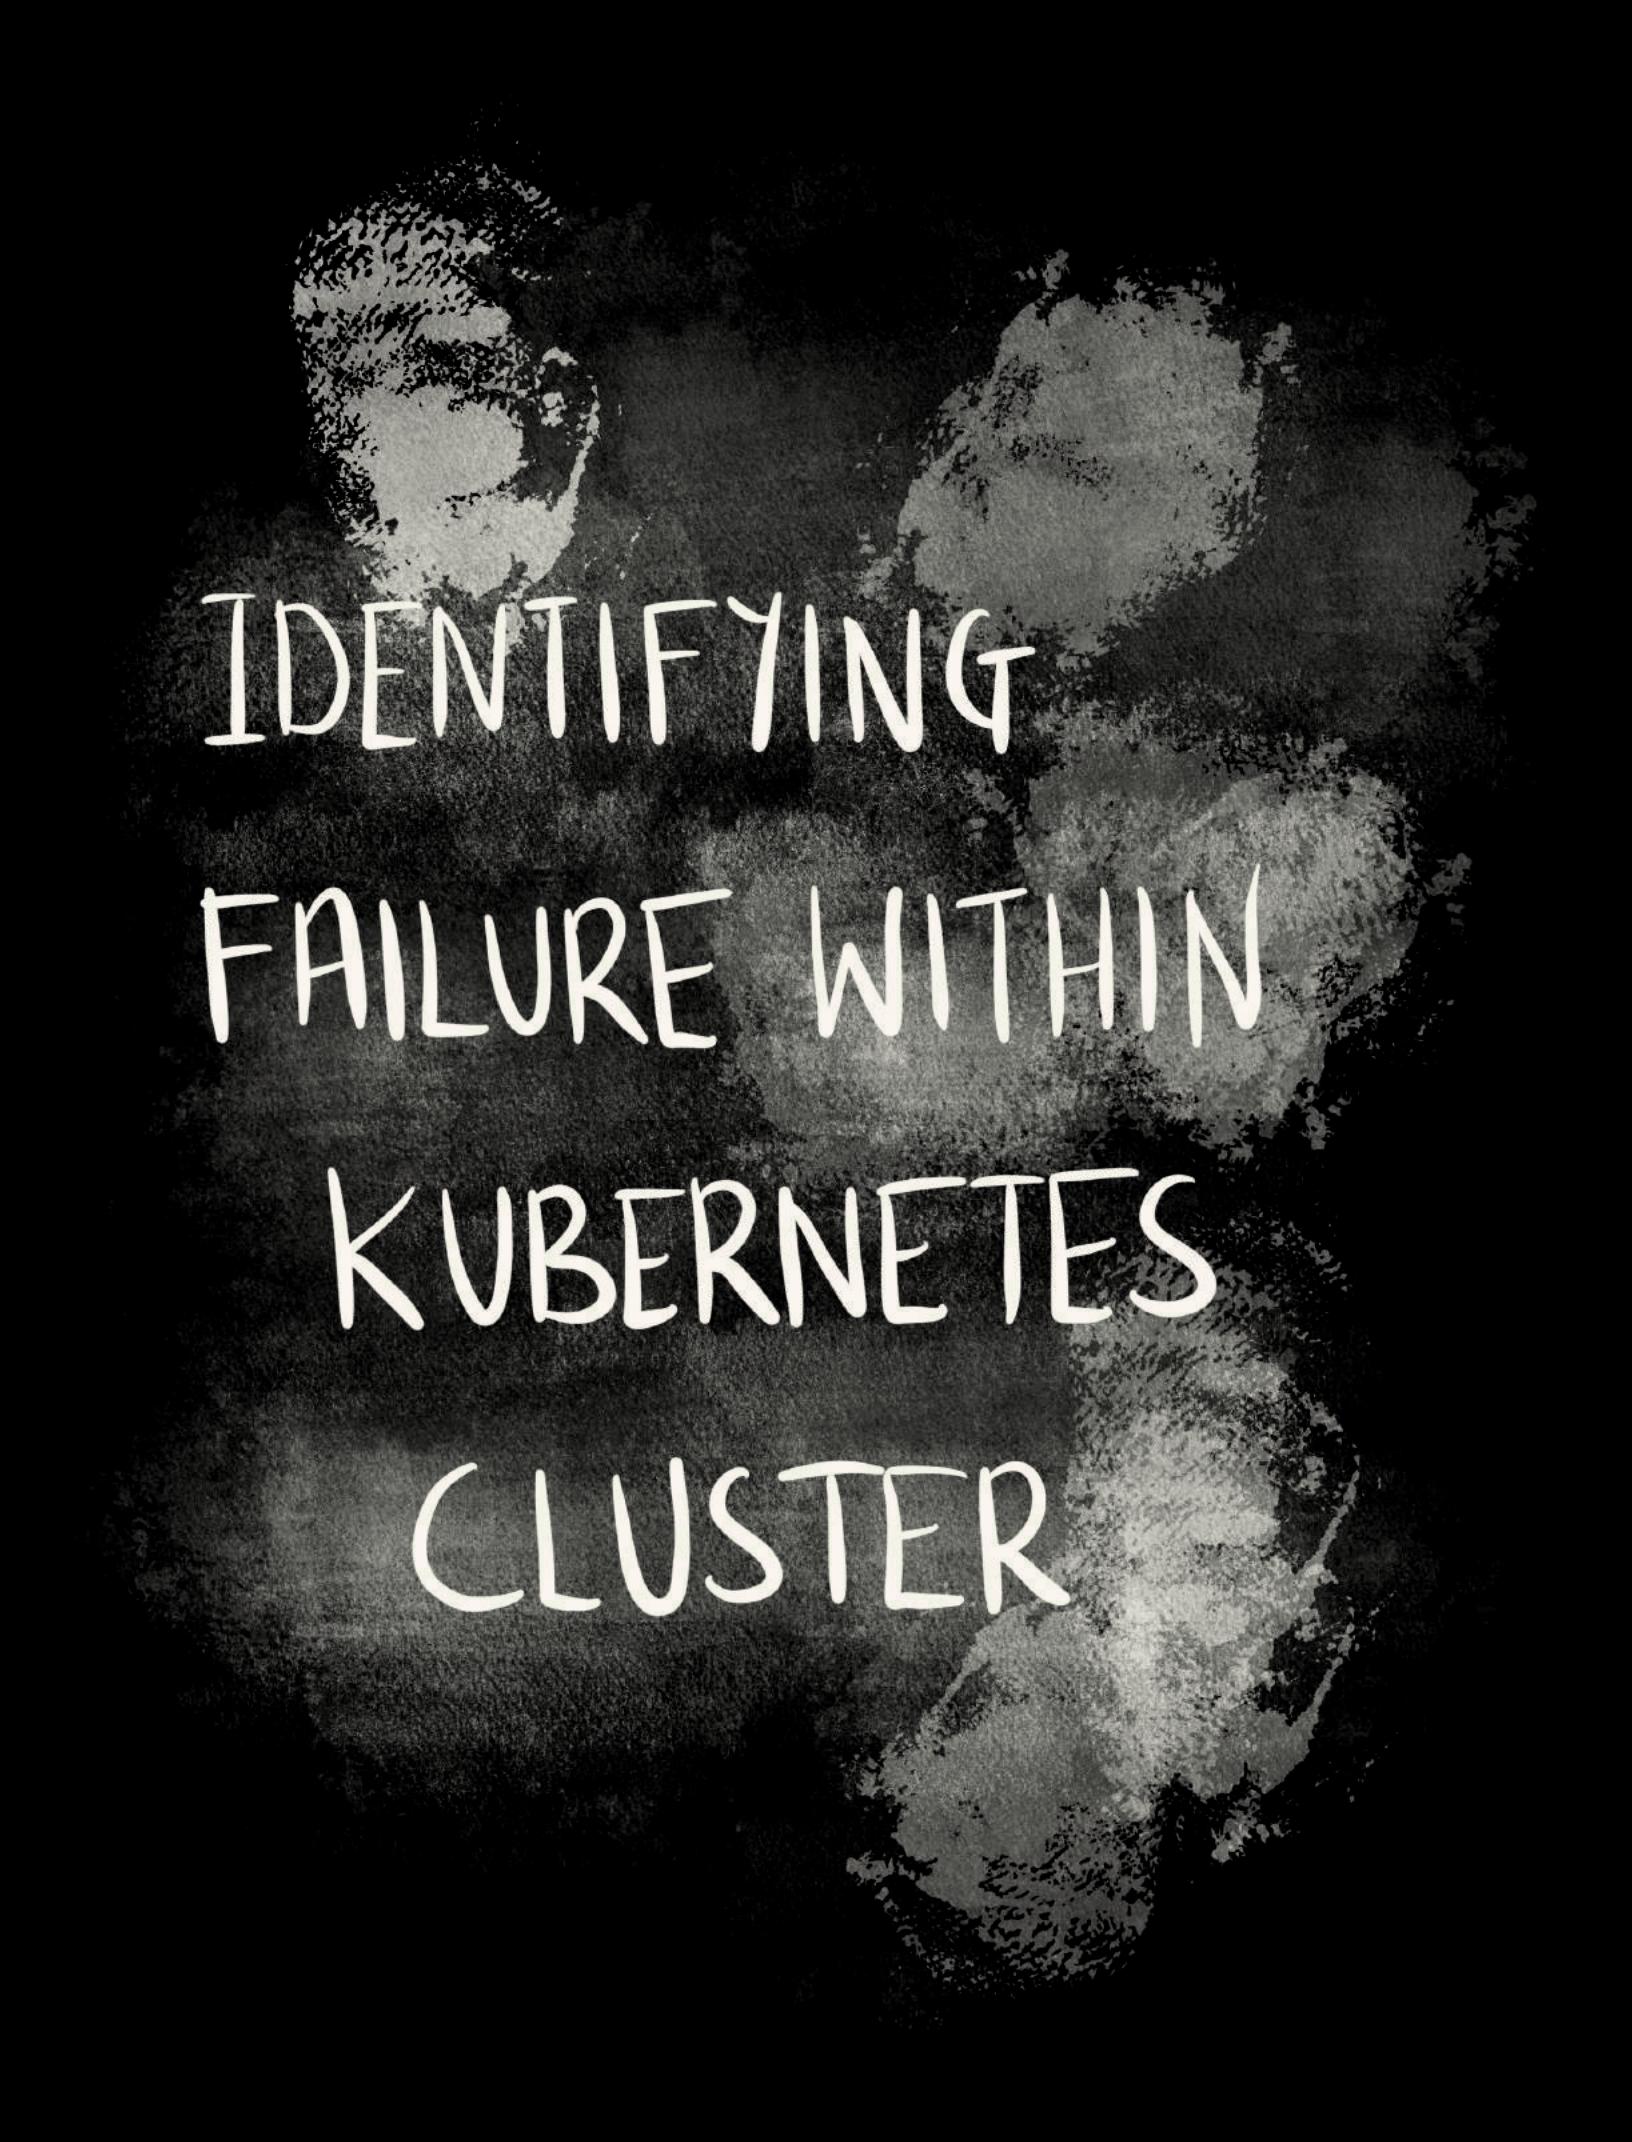
LIVENESS AND READINESS PROBE



MANAGING CLUSTER COMPONENT LOGS



- HAVE SIDEKAR CONTAINER TO DO LOGGING SO YOU CAN ACCESS SPECIFIC LOGS
- ABLE TO ROTATE LOGS USING OTHER TOOLS → NO NATIVE



IDENTIFYING FAILURE WITHIN KUBERNETES CLUSTER

TROUBLESHOOTING APPLICATION FAILURE

ABILITY TO WRITE
TERMINATION MESSAGE
TO SPECIFIC FILE ON
CONTAINER



POD1.yaml

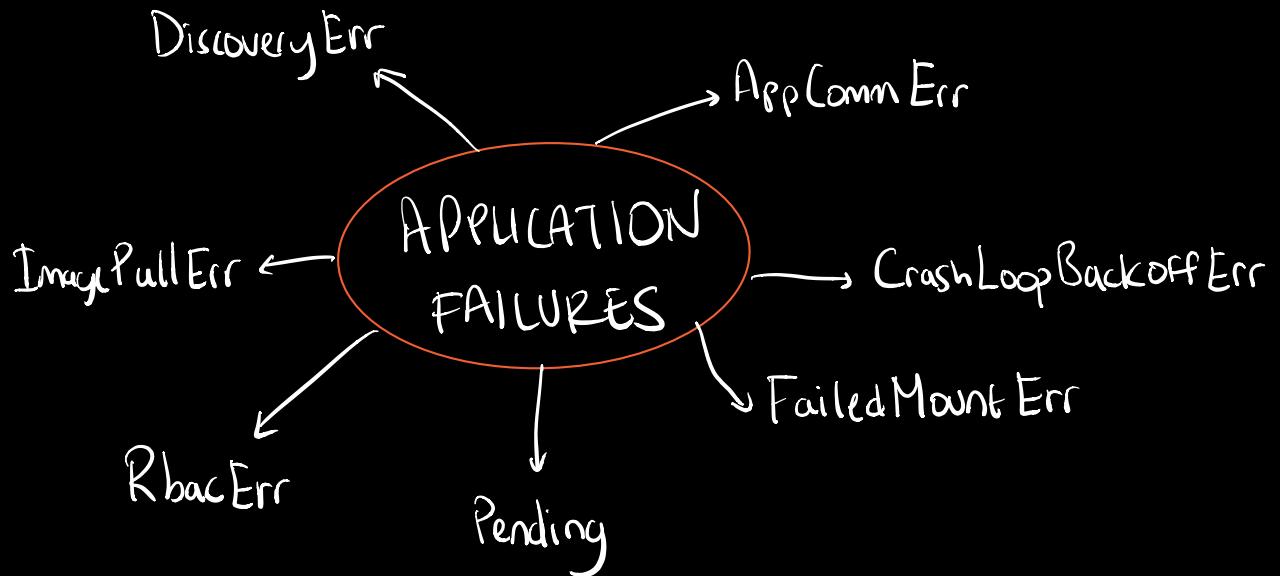
```
KIND: POD
NAME: POD1
IMAGE: busybox
COMMAND:
...
TERMINATION MESSAGE PATH
```

→ KUBECTL
DESCRIBE
↓
WILL SHOW
ERROR MESSAGE

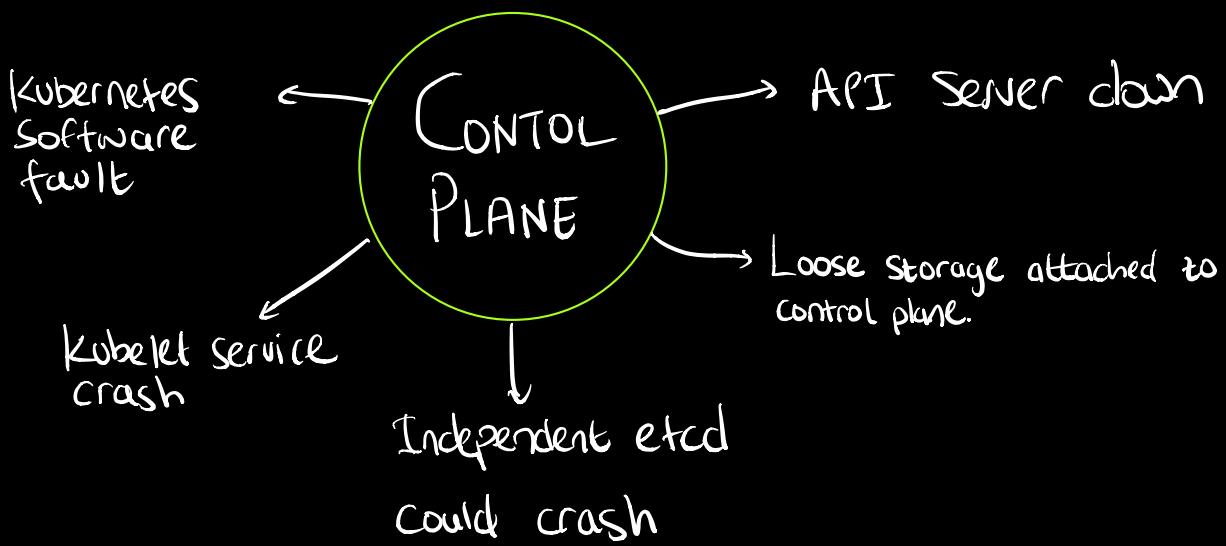
→ Only particular fields can be changed i.e Image

→ To change other fields of failed pod

↓
export configuration → -o yaml --export
modify yaml to change memory request

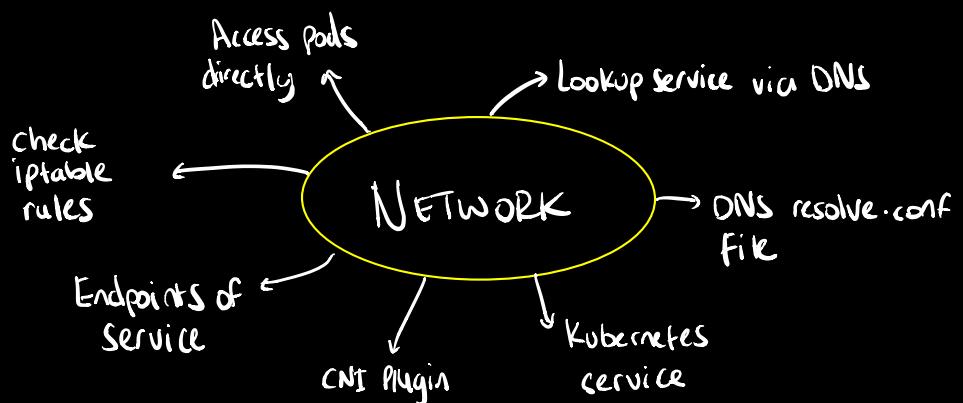
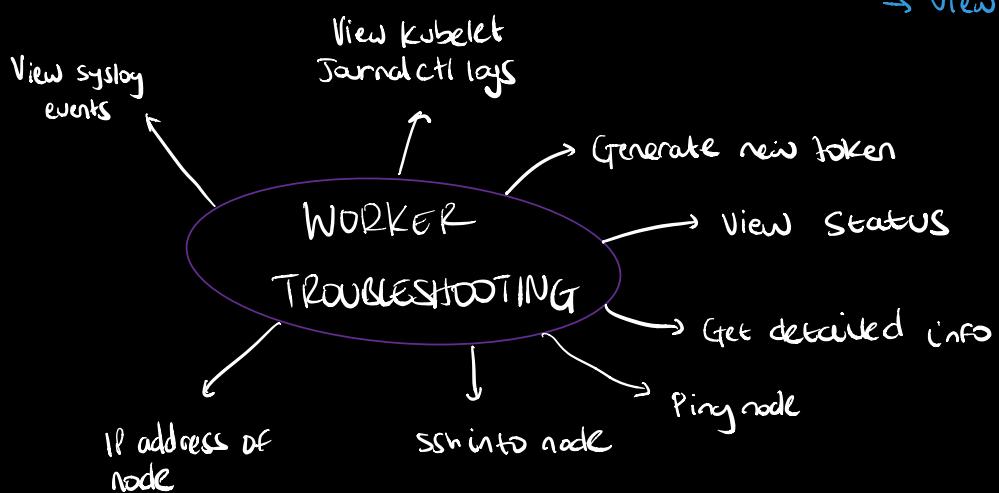


TROUBLESHOOT FAILURES



- View the events from control plane components
- View logs for control plane pods
- check status of docker service

- check status of kubelet service
- Disable swap
- check firewalld service
- View kube config



THANK You FOR READING

PLEASE LIKE AND SHARE FOR MORE

LATEST NOTES AVAILABLE ON

INSTAGRAM → adnans_techie_studies

ADNAN
RASHID
Z